

Program Efficiency Enhancement by Effective Caching

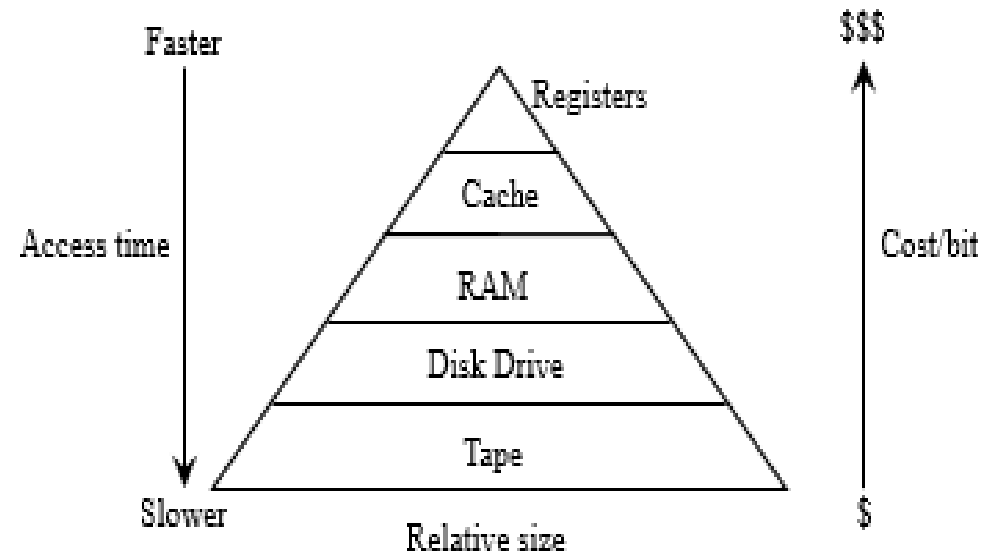
OVERVIEW OF THE C PROGRAMMING AND WRITING A BASH SCRIPT

Cache Memory

- ▶ An important part of the memory hierarchy in any computer system
- ▶ Physically located between the registers, which is part of the processor and random access memory (RAM)
- ▶ It is much larger than the number of registers and much smaller than the size of RAM
- ▶ Utilizes different technology compared to RAM
 - ▶ Faster
 - ▶ Uses more power
 - ▶ Costs more

Memory Hierarchy

- ▶ Gives the illusion of a large (RAM size), fast (cache speed) memory
- ▶ Compromise between cost, access time and size

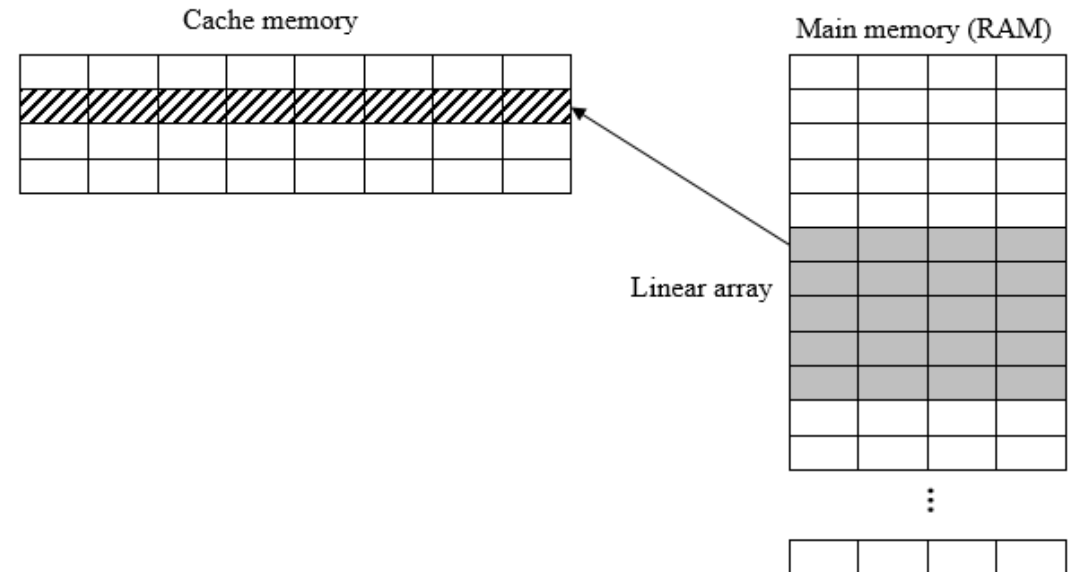


Cache Principles

- ▶ Principles of Locality
- ▶ Spatially
 - ▶ Likely nearby memory words will be accessed
 - ▶ More than one word is brought into the cache during a memory request
 - ▶ Cache line typically holds 64 bytes
- ▶ Temporally
 - ▶ Likely a memory word will be used again
 - ▶ Keep as much in the cache as possible (function of cache size)
 - ▶ Replace the least frequently used words when necessary


Efficient Memory Accesses

- ▶ Sequentially access each word in a linear array
 - ▶ First access word is not in the cache and RAM must be accessed (slow)
 - ▶ Some number of the next words are in cache (fast)
 - ▶ One slow access, $n - 1$ fast accesses



Program Overview

- ▶ The program name is **memStrideT**
- ▶ Allocates a linear array of some specified data type
- ▶ Initializes a pointer to the first element
- ▶ Iterates through all elements
 - ▶ It skips a specified number of elements called the **stride**
 - ▶ Makes multiple loops to ensure all elements are accesses
 - ▶ A value is written into each element to avoid compiler optimizations
- ▶ The time to access all array elements is output



C Language Primer

- ▶ The main function is a function that can be passed arguments
 - ▶ `int main(int argc, char* argv[])`
 - ▶ The `argc` argument means “argument count”
 - ▶ The number of arguments on the command line
 - ▶ The `argv` argument means “argument vector”
 - ▶ It is an array of pointers to characters (C string) representing the arguments value
 - ▶ If an argument is a number it needs to be converted from a **C** string into a decimal value, see the function `strtol` (string to long)

C Language Primer

- ▶ Dynamic memory allocation/deallocation
 - ▶ Use the **C** function named ***malloc***
 - ▶ It returns a pointer to a character, the first byte in the newly allocated memory block
 - ▶ The block is guaranteed to be contiguous memory locations
 - ▶ Use the **C** function ***free*** to give dynamic memory back to the Operating System
- ▶ Use the **C** function ***printf*** to output values

BASH Primer

- ▶ Linux shell can be programmed using a BASH script
- ▶ The first line in the ASCII script contains the complete pathname of the file (program) that should execute the file contents
 - ▶ **#!/bin/bash**
- ▶ Commands the would be typed can be put in the shell script
 - ▶ **make memStriderT**
- ▶ Shell scripts can define and use variables
 - ▶ **OUTFILE=memStrideDouble**
 - ▶ It is important there not be spaces

BASH Primer

- ▶ Shell scripts can contain conditional statements
- ▶ The following means remove the file if it exists in the same directory as the shell script

```
if [ -f $OUTFILE.csv ]; then
```

```
    rm -f $OUTFILE.csv
```

```
fi
```

- ▶ The value in a variable is accessed by preceding the variable name with a \$ as in \$OUTFILE above

BASH Primer

- ▶ Shell scripts can contain looping constructs

```
while [ $COUNTER -le 133 ]; do
    ./memStrideT 100 $COUNTER >> $OUTFILE.csv
    echo $COUNTER      #Progress Visualization
    let COUNTER=COUNTER+1
done
```


- ▶ This loop executes the program **memStrideT** with different values for the second argument, the number of elements to skip when accessing the linear array
- ▶ Program output is also appended to the output file during each iteration

Program Results

- ▶ The shell script takes approximately fifteen minutes to execute on a laptop with an i7 processor and 16 GB RAM
 - ▶ Surprisingly, the shell script takes even longer to execute on a massively parallel computer such as Blue Waters
- ▶ The program was executed with unsigned character, single- and double-precision floating-point elements
- ▶ In all cases, it appears the cache becomes inefficient when the stride is greater than 128 bytes
- ▶ To maximize cache efficiency it is best to access elements sequentially

Program Results

Unsigned char access time for 100 MiB



Access time for 100 MiB elements

