

In parallel processing, rather than having a single program execute tasks in a sequence, parts of the program are instead split such that the program is executed concurrently (i.e. at the same time), by multiple entities.

This implies that the execution time of a program can be made arbitrarily small by making the tasks decomposition finer in granularity.

However, concurrent tasks might need to communicate with other tasks to exchange data.

Results in communication overhead.

The tradeoff between the granularity of task decomposition and its corresponding overheads often determines the parallel performance of a program.

Parallelization Benefits

- aids in achieving speedup (aids in solving problem in less time).

- facilitates solving bigger problems.

- allows a problem that is too big to fit in the memory of one processor to be broken up such that it is able to fit in the memories of multiple processors.

Parallelization Overheads

- Interprocess interactions

 - due to communications between processors while sharing tasks data.

Idling of Processing Elements

caused due to load imbalance, synchronization, or unparallelizable serial parts of program.

Excess Computation

Best serial algorithm is difficult to parallelize; parallelizing sequential algorithm involve excess computation overhead.

Performance Metrics

It is imperative to study the performance of parallel programs

- to determine the best algorithm.

- to evaluate hardware platforms.

- to examine the benefits from parallelism.

a number of metrics are used to analyze the performance of parallel algorithms

Parallel Overhead

P = Total number of processing elements.

T_{total} = total time collectively spent by all the processing elements.

$T_{total} = P T_p$ (where T_p is time spent by a processing element).

T_s = serial time.

Total parallel overhead, T_o - time spent by all processing elements in non-useful work

$$T_o = T_{total} - T_s = P T_p - T_s$$

Speedup

The speedup of a parallel code is how much faster it runs in parallel.

If the time it takes to run a code on one processors is T_s and the time it takes to run the same code on P processors is T_p , then the speedup is given by ratio of a serial runtime to the parallel runtime

$$\text{Speedup, } S = T_s / T_p$$

T_s = Serial runtime of best sequential algorithm

Efficiency

Speedup = P (can be delivered only in ideal parallel system with P processing elements)

Processing elements cannot devote 100% time to computations.

Efficiency

a measure of how much of your available processing power is being used The simplest way to think of it is as the speedup per processor. This is equivalent to defining efficiency as the time to run P models on P processors to the time to run one model on one processor.

defined as fraction of time for which a processing element is usefully employed,

ratio of speedup to the number of processing elements

$$\text{Efficiency, } E = S/P$$

This gives a more accurate measure of the true efficiency of a parallel program than CPU usage, as it considers redundant calculations as well as idle time.

Scaling

Now that we have developed a parallel algorithm, a natural next question is, “does the algorithm scale?”

Efficiency, E can be written as $= S/P = T_s / (PT_p)$

Or

$$E = 1 / (1 + T_o / T_s)$$

Note: The total overhead function T_o is an increasing function of P .

For a given problem size

T_s remains constant

T_o increases with increase in the number of processing elements, P .

Efficiency, E of the parallel program decreases.

Amdahl's Law shows us that a program will have diminishing returns in terms of speedup as the number of processors is increased.

However, it does not place a limit on the weak scaling that can be achieved by the program, as the program may allow for bigger classes of problems to be solved as more processors become available.

The advantages of parallelism for weak scaling are summarized by John Gustafson in Gustafson's Law.

Isoefficiency Metric of Scalability

An isoefficiency function can be obtained in terms of the problem size as a function of P (no. of processing elements) to keep efficiency, E , constant.

This function determines: the ease with which a parallel system can maintain a constant efficiency.

Aids in achieving speedups in increasing proportion to the number of processing elements P .

Speedup Factors

The primary issue with speedup is the communication to computation ratio. To get a higher speedup, you can:

Maximize data locality.

Minimize volume of data exchange.

Minimize frequency of interactions.

Minimize contention and hot-spots.

Overlap computations with interactions

Replicate data or computations.

Use group communications instead of point-to-point primitives.
Overlap interactions with other interactions.