

Scalability Metrics

Parallel Processing

- ▶ In parallel processing, rather than having a **single program** execute tasks in a sequence, parts of the program are instead **split** such that the program is executed **concurrently** (i.e. at the same time), by multiple entities.
- ▶ This implies that the execution time of a program can be made **arbitrarily small** by making the tasks **decomposition finer** in granularity.
- ▶ However, **concurrent** tasks might need to **communicate** with other tasks to exchange data.
- ▶ Results in **communication overhead**.
- ▶ The **tradeoff** between the granularity of task **decomposition** and its corresponding **overheads** often determines the parallel **performance** of a program.

Parallelization Benefits

- ▶ aids in achieving **speedup** (aids in solving problem in less time).
- ▶ facilitates solving **bigger problems**.
- ▶ allows a problem that is **too big to fit in the memory** of one processor to be broken up such that it is able to fit in the memories of multiple processors.

Parallelization Overheads

- ▶ **Interprocess interactions**
 - ▶ due to communications between processors while sharing tasks data.
- ▶ **Idling of Processing Elements**
 - ▶ caused due to load imbalance, synchronization, or unparallelizable serial parts of program.
- ▶ **Excess Computation**
 - ▶ Best serial algorithm is difficult to parallelize; parallelizing sequential algorithm involve excess computation overhead.

Performance Metrics

- ▶ It is imperative to **study** the performance of **parallel programs**
 - ▶ to **determine** the best algorithm.
 - ▶ to **evaluate** hardware platforms.
 - ▶ to **examine** the benefits from parallelism.
- ▶ a **number of metrics** are used to analyze the **performance** of parallel algorithm's

Parallel Overhead

- ▶ **P**= Total number of processing elements.
- ▶ **T_{total}** = total time collectively spent by all the processing elements.
- ▶ **T_{total}** = **PT_p** (where **T_p** is time spent by a processing element).
- ▶ **T_s**= serial time.
- ▶ **Total parallel overhead, T_o**- time spent by all processing elements in non-useful work
 - ▶ **T_o**= **T_{total}** - **T_s** = **PT_p** - **T_s**

Speedup

- ▶ The **speedup** of a **parallel code** is how much **faster** it runs in parallel.
- ▶ If the time it takes to run a code on one processors is **Ts** and the time it takes to run the same code on P processors is **Tp**, then the **speedup** is given by **ratio of a serial runtime to the parallel runtime**
 - ▶ **Speedup, $S = T_s / T_p$**
- ▶ **Ts= Serial runtime of best sequential algorithm**

Efficiency



- ▶ **Speedup = P (can be delivered only in ideal parallel system with P processing elements)**
- ▶ **Processing** elements cannot **devote 100% time to computations.**
- ▶ **Efficiency**
 - ▶ a **measure** of how much of your **available processing** power is being **used** The simplest way to think of it is as the speedup per processor. This is equivalent to defining efficiency as the time to run P models on P processors to the time to run one model on one processor.
 - ▶ defined as **fraction** of **time** for which a **processing element** is usefully **employed**,
 - ▶ **ratio of speedup to the number of processing elements**
- ▶ **Efficiency, $E = S/P$**
- ▶ This gives a more **accurate measure** of the true efficiency of a **parallel program** than **CPU usage**, as it considers **redundant calculations** as well as **idle time**.

Scaling

- ▶ Now that we have developed a parallel algorithm, a natural next question is, “does the algorithm scale?”

- ▶ Efficiency, E can be written as $= S/P = T_s / (PT_p)$

- ▶ Or

- ▶ $E = 1 / (1 + T_o / T_s)$

Note: The total overhead function T_o is an increasing function of P .

- ▶ For a **given problem size**
 - ▶ T_s remains constant
 - ▶ T_o **increases** with **increase** in the number of processing elements, P .
 - ▶ Efficiency, E of the parallel program **decreases**.

Amdahl's Law

- ▶ **Amdahl's Law** shows us that a program will have **diminishing returns** in terms of speedup as the **number of processors** is **increased**.
- ▶ However, it does not place a limit on the weak scaling that can be achieved by the program, as the program may allow for bigger classes of problems to be solved as more processors become available.
- ▶ The **advantages of parallelism for weak scaling** are summarized by John Gustafson in Gustafson's Law.

Isoefficiency Metric of Scalability

- ▶ An **isoefficiency function** can be obtained in terms of the problem size as a function of **P (no. of processing elements)** to keep **efficiency, E**, constant.
- ▶ **This function determines:** the ease with which a parallel system can maintain a constant efficiency.
- ▶ Aids in **achieving speedups** in increasing **proportion** to the number of processing elements **P**.

Speedup Factors

- ▶ **The primary issue with speedup is the communication to computation ratio. To get a higher speedup, you can:**
- ▶ Maximize data locality.
- ▶ Minimize volume of data exchange.
- ▶ Minimize frequency of interactions.
- ▶ Minimize contention and hot-spots.

Speedup Factors

- ▶ Overlap computations with interactions
- ▶ Replicate data or computations.
- ▶ Use group communications instead of point-to-point primitives.
- ▶ Overlap interactions with other interactions.

Acknowledgements

Adapted from:

- ▶ Chapter 3 slides, by A. Grama, of the text: Principles of Parallel Algorithm Design by Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar

Thank You!

