

Markov Chains

SERIAL AND PARALLEL MATRIX MULTIPLICATION

Learning Objectives


- ▶ The student should understand
 - ▶ The concept of a Markov Chain
 - ▶ The transition matrix
 - ▶ State changes using the transition matrix and a state vector (matrix-vector multiplication)
 - ▶ The steady state (iterated matrix-matrix multiplication)
 - ▶ How simple it **can be** to convert a serial program into a parallel program using OpenMP

Markov Chain

- ▶ A mathematical representation of a problem involving state transitions
- ▶ The transitions are based on a fixed probability
 - ▶ When the system is in one state the probability of moving to another state is fixed
- ▶ Useful to visualize a Markov chain with a state space diagram
 - ▶ Each state is represented as a node (vertex or circle)
 - ▶ Probabilities of transitioning from one state to another state
 - ▶ Weighted edges between the nodes

Markov Chain Example

- ▶ Three vats of water, A , B , and C
 - ▶ 10% of the water in A is transferred into B and 20% in A is transferred into C in one hour
 - ▶ 20% of the water in B is transferred into A and 20% from B is transferred into C in one hour
 - ▶ 30% of the water in C is transferred into A and 10% from C is transferred into B in one hour
- ▶ The state-space diagram



Markov Chain Example

- ▶ The question to be answered is, how much water is in each vat after a certain number of hours
- ▶ The state space diagram can be represented as a two-dimensional matrix, called a transition matrix
 - ▶ The columns represent the transition amounts, or the probabilities
 - ▶ The columns must sum to one

$$T = \begin{bmatrix} 0.7 & 0.2 & 0.3 \\ 0.1 & 0.6 & 0.1 \\ 0.2 & 0.2 & 0.6 \end{bmatrix}$$

- ▶ $T[r][c]$ represents the percentage of water that is taken from vat c and transferred to vat r

Markov Chain Example

- ▶ A vector is used to represent the initial amount of water in each vat
- ▶ For example, let vat A initially contain 5 gallons of water, vat B contain 10 gallons, and vat C contain 15 gallons

$$V_0 = \begin{bmatrix} 5 \\ 10 \\ 15 \end{bmatrix}$$

- ▶ The amount of water in each vat is determined by multiplying the transition matrix by the vector, or

$$V_1 = TV_0$$

$$V_1 = TV_0 = \begin{bmatrix} 0.7 & 0.2 & 0.3 \\ 0.1 & 0.6 & 0.1 \\ 0.2 & 0.2 & 0.6 \end{bmatrix} \begin{bmatrix} 5 \\ 10 \\ 15 \end{bmatrix} = \begin{bmatrix} 10 \\ 8 \\ 12 \end{bmatrix}$$

Markov Chain Example

- ▶ The previous equation states that after one hour vat A will contain 10 gallons, vat B will contain 8 gallons, and vat C will contain 12 gallons
- ▶ The amount of water in each vat after the next hour is found in a similar manner, using these new quantities as the initial conditions

$$V_2 = TV_1 = T[TV_0] = \begin{bmatrix} 0.7 & 0.2 & 0.3 \\ 0.1 & 0.6 & 0.1 \\ 0.2 & 0.2 & 0.6 \end{bmatrix} \begin{bmatrix} 10 \\ 8 \\ 12 \end{bmatrix} = \begin{bmatrix} 12.2 \\ 7 \\ 10.8 \end{bmatrix}$$

- ▶ Likewise, the amount of water in each vat after N hours is

$$V_n = T \dots [TV_0] = T \dots TV_0 = T^N V_0$$

Markov Chain Example

- ▶ The previous equation demonstrates it is not necessary to compute each state to determine the state after N hours, but rather you need to multiply the transition matrix by itself N times

Matrix Multiplication

- ▶ When multiplying two matrices together, the number of columns of the matrix on the left must be the same as the number of rows of the matrix on the right
- ▶ The final result of multiplying two matrices together is a third matrix, and the process can be described by the equation,

$$C = AB$$

- ▶ Where A, B, and C are matrices
- ▶ Each element in the product matrix is formed by multiplying all the elements of a row with the corresponding elements of a column
- ▶ For example,

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}, B = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$$

Matrix Multiplication

- ▶ Then

$$C = \begin{bmatrix} 1(5) + 2(7) & 1(6) + 2(8) \\ 3(5) + 4(7) & 3(6) + 4(8) \end{bmatrix} = \begin{bmatrix} 19 & 22 \\ 43 & 50 \end{bmatrix}$$

- ▶ Which can be written mathematically as

$$c_{i,j} = \sum_{k=0}^{n-1} a_{i,k} b_{k,j}$$

- ▶ Each element in the C matrix requires n multiplications
- ▶ Matrix multiplication is an n^3 operation, which is prohibitively expensive for large matrices

Matrix Multiplication

- ▶ The code to perform this computation is straightforward
 - ▶ For all rows and all columns, compute each element in the resulting matrix
 - ▶ This code assumes the matrix C is initialized with zeros

```
for (int r = 0; r < n; r++)  
    for (int c = 0; c < n; c++)  
        for (int k = 0; k < n; k++)  
            C[r][c] += A[r][k] * B[k][c];
```

Parallel Matrix Multiplication


- ▶ Matrix multiplication is an example of an embarrassingly parallel problem
 - ▶ All computations are independent and can therefore be performed simultaneously
- ▶ The power of OpenMP is to simply convert serial code for an embarrassingly parallel problem into a parallel solution by adding a single directive to the serial code

```
#pragma omp parallel for shared(A, B, C, n)
    for (int r = 0; r < n; r++)
        for (int c = 0; c < n; c++)
            for (int k = 0; k < n; k++)
                C[r][c] += A[r][k] * B[k][c];
```

Parallel Matrix Multiplication

- ▶ The pragma tells the compiler to break up the A matrix into a number of blocks containing a similar number of rows, and have independent threads perform the computation on each block of rows simultaneously
 - ▶ Each thread performs fewer computations compared to the serial case, therefore improving the performance of the matrix multiply code
 - ▶ There is some overhead of performing the computation in parallel
 - ▶ The speedup is less than one over the number of threads
 - ▶ The speedup becomes more significant as the size of the matrix increases
 - ▶ The number of hardware threads also limits any performance improvement
- ▶ The following plot shows the time to perform matrix multiplication on various sized matrices using different numbers of threads

Parallel Matrix Multiplication Efficiency



Conclusion

- ▶ Matrix multiplication is the fundamental processing step when a problem can be mathematically modeled with a Markov Chain
- ▶ Matrix multiplication is an embarrassingly parallel computational problem that can be easily made parallel with OpenMP
- ▶ Dramatic improvements in computation time can be observed with a parallel version of matrix multiplication, especially as the size of the matrices get large