

# OpenMP Tasks


Cameron Foss

Blue Waters Capstone

# Outline

- What are OpenMP tasks?

- task constructs
- directives



```
void recurse()
{
    ... ..
    recurse();
    ... ..
}

int main()
{
    ... ..
    recurse();
    ... ..
}
```

recursive call


- How to implement in Code

- “A race car” → a hello-world to OpenMP tasking
- OpenMP barrier and taskwait
- Fibonacci recursive example

1,2,3,5,8,13 ...


- Assignment using recursion

- Compute pi recursion: implement omp task directives



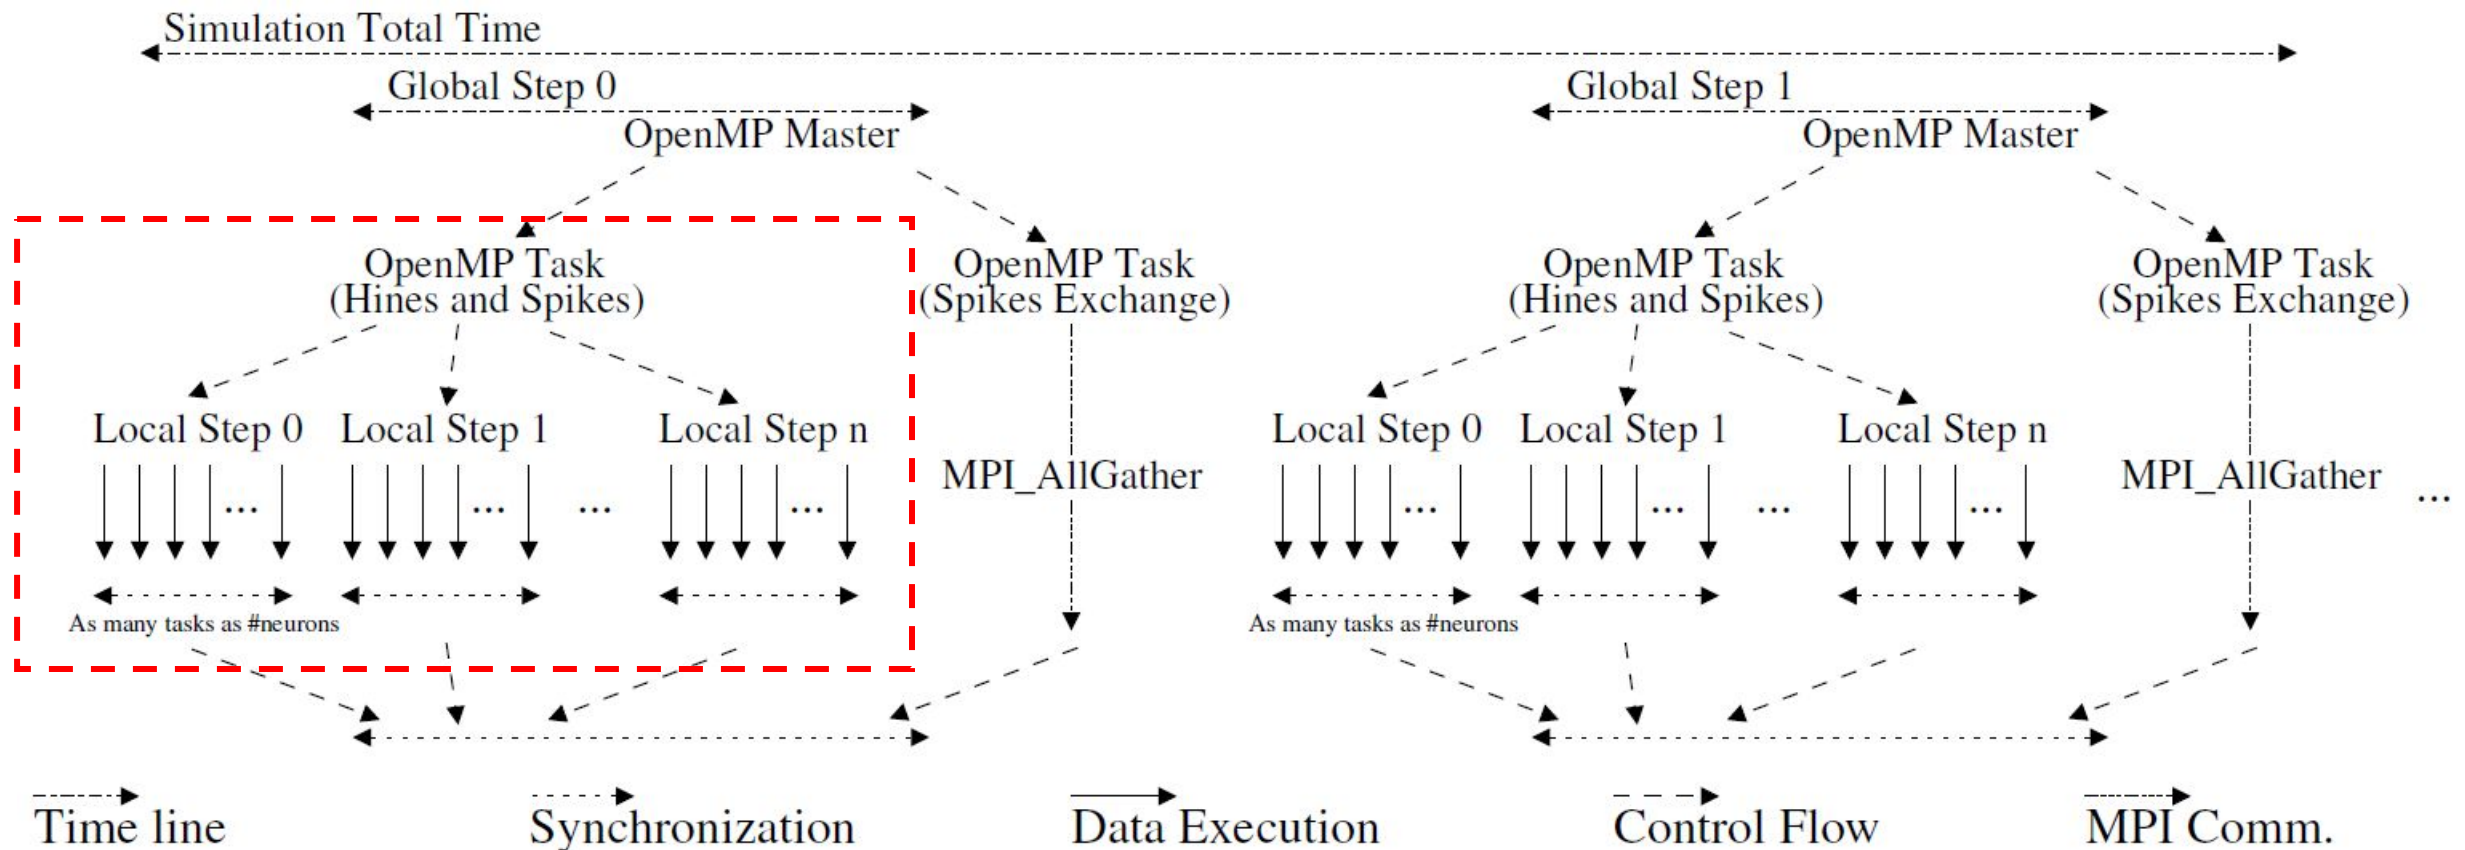
# OpenMP Tasking Model

*Tasking facilitates the parallelization of applications where units of work are generated dynamically, as in **recursive** structures or **while** loops.*  
[API User's Guide]



# A more detailed example...

Parallel model implemented for the simulator based on MPI+OpenMP tasking.



# Why use tasking...?

- Up until OpenMP 3.0, certain types of parallelism was not feasible
  - This required messy workarounds at best
- Threads all run the same code. Tasks allow us to assign different blocks of code that are independent to different threads.
- enable irregular parallelism
- Allowed specific code blocks to be run simultaneously and independently of other task blocks.

# Task Execution

- Tasks are generated by specifying the task construct.
- Tasks can be generated anywhere in the code.
- A task is composed of:
  - Code to be executed
  - Data environment (inputs to be used and outputs to be generated)
  - A location where the task will be executed (a thread)

# Task Types

- **explicit** tasks are generated using the task directive

```
# pragma omp task clause
{
    block of code ...
}
```

- **implicit** tasks are inherent to parallel constructs and function similarly to explicit tasks

# Data environment

*clause* can be:

- **depend (*list*)**
- **if (*expression*)**
- **final (*expression*)**
- untied
- mergeable
- **default (*shared* | *firstprivate* | *none*)**
- **private (*list*)**
- **firstprivate (*list*)**
- **shared (*list*)**
- priority (*value*)

defines storage locations  
where task dependencies exist

Generate a task if  
a condition is met

Stop generating new  
tasks if a condition is met

Memory storage requirements  
on input output variables

```
# pragma omp task clause
{
    block of code ...
}
```

- inputs to be used and outputs to be generated
- The task directive takes the following data-sharing attribute clauses that define the data environment of the task.




# Tasking Explained...

**Objective:** Write a program that prints either “A race car” or “A car race” and maximize the parallelism.



# A race car OR A car race...(1)



```
#include <stdlib.h>
#include <stdio.h>


int main(int argc, char *argv[])
{
    printf("A ");
    printf("race ");
    printf("car ");

    printf(".\n");
    return(0);
}
```

Q: What will this program print?

```
$ gcc racecar.c
$ ./a.out
  A race car .
$
```

# A race car OR A car race...(2)



```
#include <stdlib.h>
#include <stdio.h>


int main(int argc, char *argv[])
{
    #pragma omp parallel
    {
        printf("A ");
        printf("race ");
        printf("car ");
    } // End of omp parallel region

    printf(".\n");
    return(0);
}
```

Q: What will this program print when we use 2 threads?

```
$ gcc -fopenmp racecar_omp.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A race car A race car .
$
```

# A race car OR A car race...(3)



```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            printf("race ");
            printf("car ");
        }
    } // End of omp parallel region

    printf(".\n");
    return(0);
}
```

Q: What will this program print when we use 2 threads?

```
$ gcc -fopenmp racecar_omp.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A race car .
$
```

**It only prints it once now because the omp single block is run by only one (*single*) thread.**

# A race car OR A car race...(4)



```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            #pragma omp task
            { printf("race "); }
            #pragma omp task
            { printf("car "); }
        }
    } // End of omp parallel region

    printf(".\n");
    return(0);
}
```

Q: What will this program print when we use 2 threads?

```
$ gcc -fopenmp racecar_omp.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A race car .
$ ./a.out
A car race .
```

**Tasks are  
executed in  
arbitrary order!**

Note: It may take many attempts before both "A race car" and "A car race" appear.

# A race car OR A car race...(5)

**Objective:** Write a program that prints either “A race car” or “A car race” and maximize the parallelism.

**COMPLETED**

**Objective 2:** Add “is fun to watch.” to the end of the sentence...

# A race car OR A car race...(6)



```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            #pragma omp task
            { printf("race "); }
            #pragma omp task
            { printf("car "); }
            printf("is fun to watch ");
        }
    } // End of omp parallel region


    printf(".\n");
    return(0);
}
```

Q: What will this program print when we use 2 threads?

```
$ gcc -fopenmp racecar_omp.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A is fun to watch race car .
$ ./a.out
A race is fun to watch car .
...
```



# A race car OR A car race...(7)



```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            #pragma omp task
            { printf("race "); }
            #pragma omp task
            { printf("car "); }
            #pragma omp taskwait
            printf("is fun to watch ");
        }
    } // End of omp parallel region
    printf(".\n");
    return(0);
}
```


Task 1 is generated

Task 2 is generated

**Synchronization point:**  
"Wait here until all  
previous tasks have  
finished..."



# A race car OR A car race...(8)



```
#include <stdlib.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            printf("A ");
            #pragma omp task
            { printf("race "); }
            #pragma omp task
            { printf("car "); }
            #pragma omp taskwait
            printf("is fun to watch ");
        }
    } // End of omp parallel region
    printf(".\n");
    return(0);
}
```

Q: What will this program print when we use 2 threads?

```
$ gcc -fopenmp racecar_omp.c
$ export OMP_NUM_THREADS=2
$ ./a.out
A race car is fun to watch .
$ ./a.out
A car race is fun to watch .
...
```

# A race car OR A car race...(9)

**Objective:** Write a program that prints either “A race car” or “A car race” and maximize the parallelism.

**COMPLETED**

**Objective 2:** Add “is fun to watch.” to the end of the sentence... **COMPLETED**

# Recap so far...

- When a thread encounters a `#pragma omp task` it generates a task
  - That is, the number of threads that encounter a `#pragma omp task` will generate that many number of tasks.
  - Can control with a directive `#pragma omp single` or more dynamically with the function `omp_set_num_threads`
- Tasks are executed in parallel and in **arbitrary order**.
  - Implies a race occurs between tasks that are initiated by the same thread.
- `#pragma omp taskwait` can be used to execute code by the main thread strictly after all tasks are completed.

# Recursion

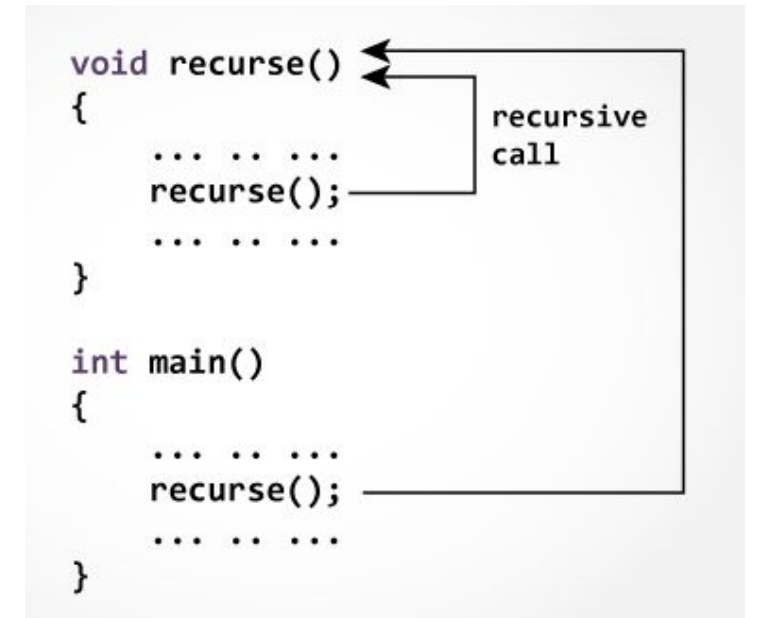
- Tasking is often used in codes that use recursion
- **Recursion**, in simple terms, means defining a problem in terms of itself.
- A **recursive function** is a function that calls itself iteratively until some condition is met.

**Recursion** is defined as “relating to or involving a program or routine of which a part requires the application of the whole, so that its explicit interpretation requires in general many successive executions.”

Calculating the  $n^{th}$  element in a Fibonacci sequence

$\text{fib}(n = 6) = [1, 2, 3, 5, 8, \mathbf{13}]$   
 $= 13$

```
int fib(int n) {  
    int i, j;  
    if (n<2) { return n; }  
    else {  
        i=fib(n-1);  
        j=fib(n-2);  
        return i+j;  
    }  
}
```



# Recursion – Compute Pi

```
double pi_r (double h, unsigned depth, unsigned maxdepth, unsigned long long begin, unsigned long long niters)
{
    if (depth < maxdepth)
    {
        double area1, area2;
        // Process first half
        area1 = pi_r (h, depth+1, maxdepth, begin, niters/2-1);
        // Process second half
        area2 = pi_r (h, depth+1, maxdepth, begin+niters/2, niters/2);
        return area1+area2;
    }
    else
    {
        unsigned long long i;
        double area = 0.0;
        for (i = begin; i <= begin+niters; i++)
        {
            double x = h * (i - 0.5);
            area += (4.0 / (1.0 + x*x));
        }
        return area;
    }
}
```

# Recursion – Compute Pi

```
double pi (unsigned long long niters)
{
    double res;
    double h = 1.0 / (double) niters;
#define MAX_PARALLEL_RECURSIVE_LEVEL 4
    res = pi_r (h, 0, MAX_PARALLEL_RECURSIVE_LEVEL, 1, niters);
    return res * h;
}

int main (int argc, char *argv[])
{
    double tstart = omp_get_wtime();
    int NITERS = atoi(argv[1]);
    printf ("Number of rectangles: %d\n",NITERS);

    double pi_approx = pi(NITERS);


    printf ("PI (w/%d iters) is %15.13f\n", NITERS, pi_approx);
    printf ("Accuracy for %d \titers is :: %10.7f % \n", NITERS, pi_approx/M_PI*100);
    printf ("runtime is: %lf\n", omp_get_wtime()-tstart);
    printf ("\n");
    return 0;
}
```

# Assignment: Compute pi with OpenMP tasks

1. Parallelize the compute pi code with openMP
2. Use openMP tasks on the recursive function pi\_r
3. Once the code has been parallelized, perform the following:
  - a) Compute pi with [100,500,1000,5000,10000,1000000,100000000] numbers of rectangles for both the serial and parallelized code with 2 and 4 threads. Note the runtime and accuracy of the approximation in each.
  - b) Comment on trends you see with the runtime when increasing the number of rectangles and the number of threads.

$$\frac{\lim_{n \rightarrow \infty} \sum_n h[n] \times dx}{r^2} \approx \pi$$

$$h[n] = (r^2 - x[n]^2)^{1/2}$$



# References and Additional resources:

## OpenMP Tasking :

- <https://openmp.org/wp-content/uploads/sc13.tasking.ruud.pdf>
- <http://icl.cs.utk.edu/classes/cosc462/2017/pdf/W43%20-%20OpenMP%20Tasking.pdf>
- <https://en.wikibooks.org/wiki/OpenMP/Tasks>

## OpenMP tasking model (advanced):

- [https://www.slideshare.net/InformaticaUCM/openmp-tasking-model-from-the-standard-to-the-classroom?from\\_action=save](https://www.slideshare.net/InformaticaUCM/openmp-tasking-model-from-the-standard-to-the-classroom?from_action=save)

## A nice book on OpenMP in general, many many examples:

- <https://www.openmp.org/wp-content/uploads/OpenMP4.0.0.Examples.pdf>

## OpenMP API User's Guide:

- [https://docs.oracle.com/cd/E77782\\_01/html/E77801/gliyr.html#scrolltoc](https://docs.oracle.com/cd/E77782_01/html/E77801/gliyr.html#scrolltoc) (2017 documentation)
- <https://docs.oracle.com/cd/E19205-01/820-7883/auto15/index.html> (2010 documentation)

## OpenMP directives:

- <https://docs.microsoft.com/en-us/cpp/parallel/openmp/reference/openmp-directives?view=vs-2019>

## Basics of recursion

- <https://www.programiz.com/c-programming/c-recursion>
- <https://www.cs.utah.edu/~germain/PPS/Topics/recursion.html>

## Runtime Determinacy Race Detection for OpenMP Tasks:

- [https://link.springer.com/chapter/10.1007/978-3-319-96983-1\\_3](https://link.springer.com/chapter/10.1007/978-3-319-96983-1_3)