



Acceleration of Stencil Code Using MPI

A Hands-on Approach

Learning Objectives

- **APPLY** basic MPI functions to accelerate a common scientific pattern
- **USE** Profiler tools to evaluate bottle necks in code
- **IDENTIFY** which loops/task can be distributed (loop dependency)
- **COMPARE** sequential to distributed execution times

A large orange circle is positioned on the left side of the slide, partially cut off by the edge.

Expected Time of the Activity:

Module Approximate Timing

- Problem Introduction : 5 min
- Sequential Code: 5 min
- Profiling and Identification of main bottleneck, stencil loop: 2 min
- Explain how to accelerate this loop using MPI : 12min
- Compare Sequential vs Parallel OpenMP implementation: 1 min

Total time for the module: 25 minutes

A series of yellow dashed line segments are arranged in a curved path at the bottom right of the slide.

Problem Introduction : Stencil Code

Stencil code: are a class of iterative functions which update an array according to some pattern (stencil)

They are most found in:

- ☐ Scientific Simulations:
(ex. Heat Transfer)
- ☐ Computational Fluid Dynamics
(ex. PDE solvers, Poisson Equation)
- ☐ Image Processing
(ex. Image filters)
- ☐ Machine Learning
(ex. Convolutional Neural Networks)





Figure 1 Stencil code on a 2 D array, using 9 neighbors to calculate the output

Example Sequential 4 neighbors Code in C


No Boundary Consideration

```
for (int i = 0; i < iter_count; ++i) {  
    for (int y = 0; y < N; y++) {  
        for (int x = 0; x < N; x++) {  
            //ctop,cbottom,ceast,cright are the coefficients  
            //of the stencil or filter  
            out[y][x]=in[y][x]+(ctop*in[y-1][x] +  
                                cbottom*in[y+1][x] +  
                                cwest*in[y][x-1] +  
                                ceast*in[y][x+1] )/SPEED;  
        }  
    }  
    //swap input and output array  
    tmp= out; out=in; in=tmp;  
}
```



Example Sequential 4 neighbors Code in C With Boundary

```
for (int i = 0; i < iter_count; ++i) {  
    for (int y = 0; y < height; y++) {  
        for (int x = 0; x < width; x++) {  
            //o[y][x] is same as o[y*N+x]  
            center=y*N+x;  
            west= (x==0) ?center:center-1; east= (x==N-1) ?center:center+1;  
            top= (y==0) ?center:center-N; bottom= (y==N-1) ?center:center+N;  
            out[y*N+x]=in[center]+ (ctop*in[top] +  
                                    cbottom*in[bottom] +  
                                    cwest*in[west] +  
                                    ceast*in[east] )/SPEED;  
        }  
    }  
    //swap input and output array  
    tmp= out; out=in; in=tmp;  
}
```



NOTE :Explaining Boundary condition notation

If (row x is equal to 0) then west=center

else west=center-1

Can be written in C as one compound statement

west= (x==0) ?center:center-1

Sequential Code Profiling

Full code can be found in file: stencil.c
Compile : gcc -o stencil stencil.c
CPU: Intel i7-7820HQ 8 cores @2.90GHz
gcc -version: 9.3.0

Elapsed time 831.090 sec. (for an array of size 2^{11} by 2^{11} ,
6553 iterations(convolutions) and speed 10^8)

Use

Use [gprof](#) or [perf](#) (profilers) to analyze timings and bottleneck
on code:



Function	Execution %
Init function	0.1%
Stencil	99.9%
PrintResult	0%

in our case is very clear the function that needs acceleration is
the Stencil for loop shown on previous slide (6)

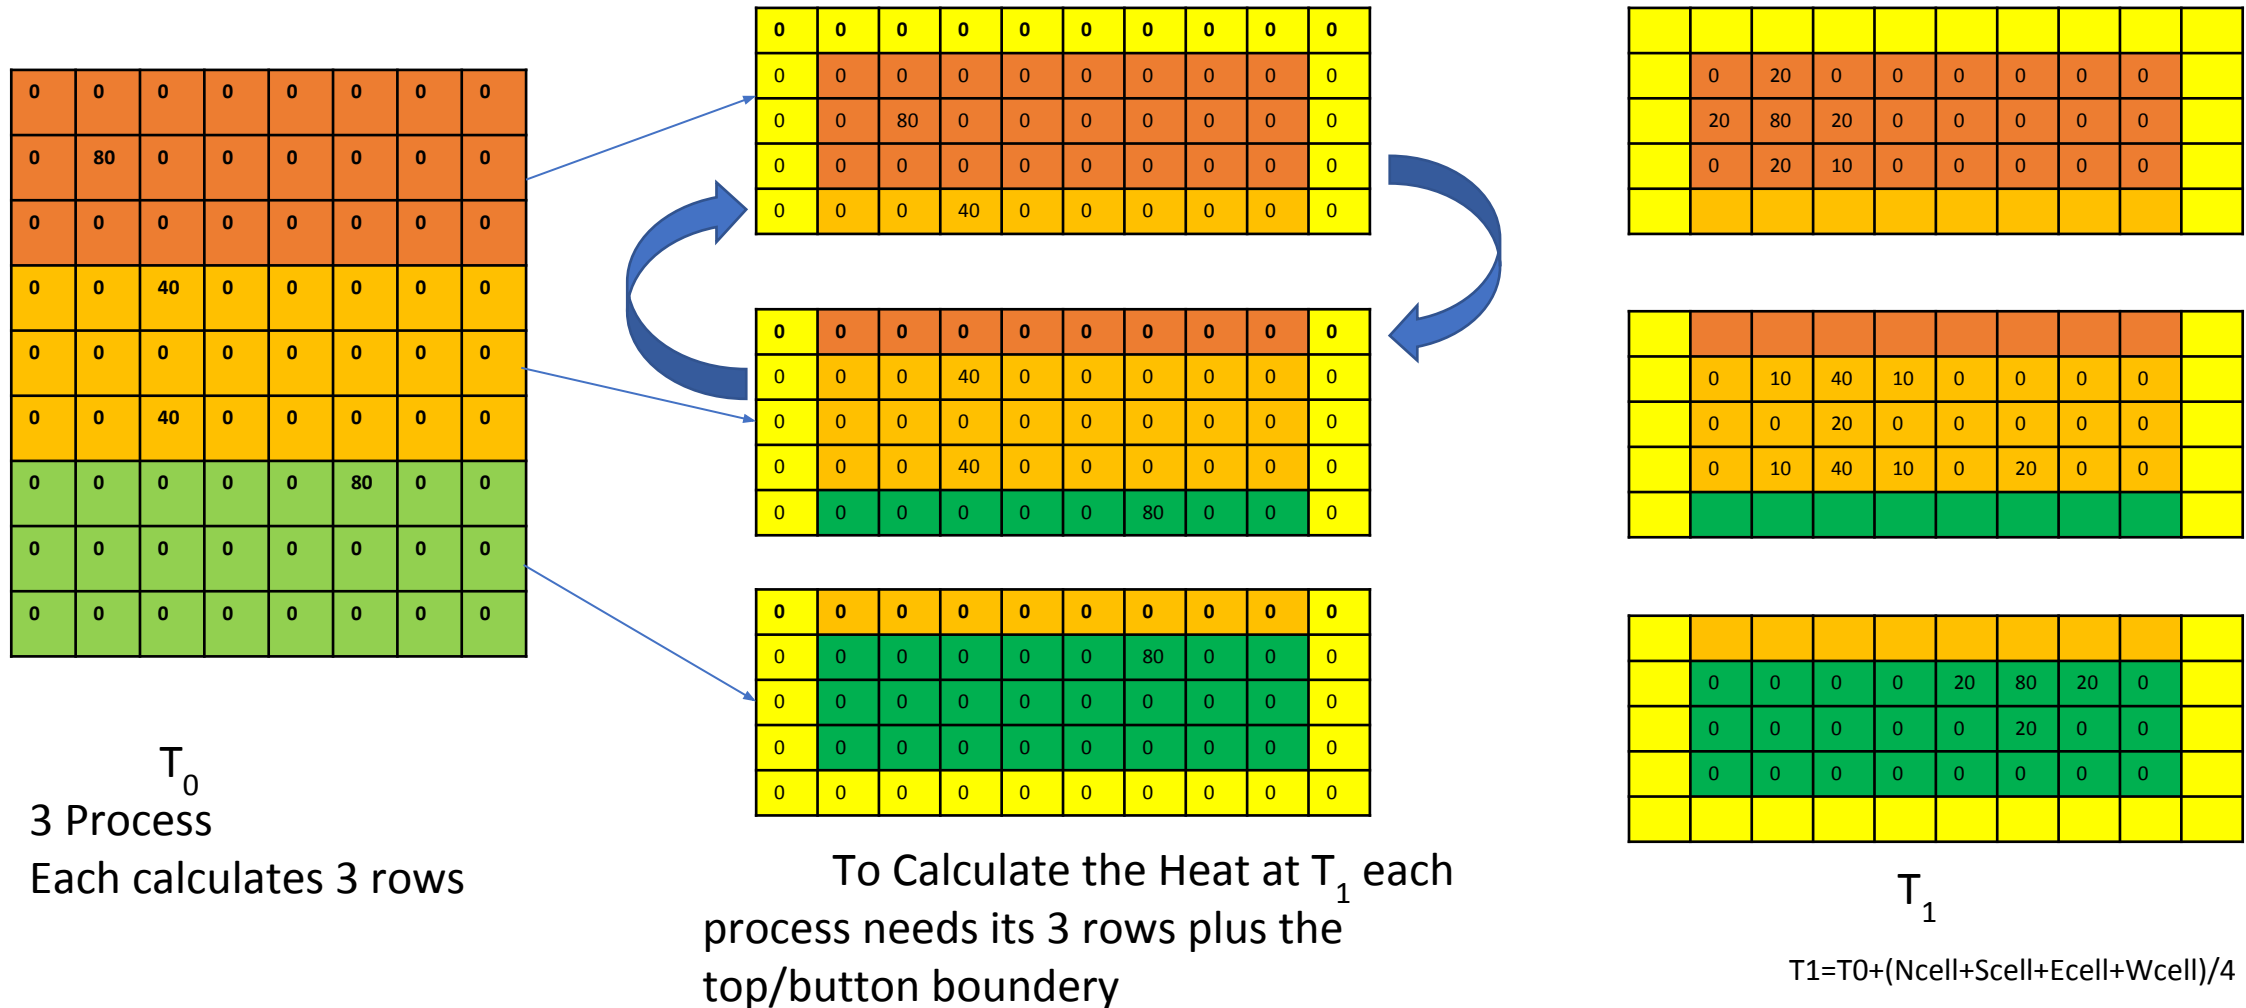


Output Visualization



- The Input to the program is a 2D array , each cell in the array contains the initial temperature
- The Output of the program also a 2D array , contains the temperatures after an X amount of time
- The 2D arrays output can be displayed as png images




MPI Acceleration of main loop. Divide Array in rows



MPI Acceleration of main loop. Divide Array in rows


$$T_1$$


exchange


$$T_2$$

MPI Stencil

//Distribute the Data Matrix among the processes

```
int myRows=ROWS/numProcess;
```

```
float *myData=new float[myRows*cols];
```

```
float *buff= new float[myRows*cols] ;
```

```
MPI_Scatter(data, myRows*cols, MPI_FLOAT, myData, myRows*cols,MPI_FLOAT);
```

```
memcpy(buff,myData, myRows*cols*sizeof(float));
```

//exchange the boundary rows

```
for (int iter=0; iter<numIter;iter++){
```

```
    if(myID>0){
```

```
        MPI_Isend(myData,cols,MPI_Float,myID-1,0); //send first row of myData to previous process
```

```
        MPI_Irecv(prevRow,cols,MPI_Float,myID-1,0); //receives last row from previous process
```

```
    }
```

```
    if(myID<numProcess-1){
```

```
        MPI_Isend(&myData[(myRows-1)*cols],cols,MPI_Float,myID+1,0); //send last row of myData to next process
```

```
        MPI_Irecv(nextRow,cols,MPI_Float,myID+1,0); //receives first row from the next process
```

```
    }
```

MPI Stencil

```
//calculate main Loop
```

```
for(int i=1; i<myRows-1;i++){  
    for(int j=1; j<cols-1;j++){  
        buff[i*cols+j] = myData[i*cols+j]+( cbotton* myData[(i+1)*cols+j]+  
                                              cwest*cmyData[i*cols+j-1] +  
                                              ceast*myData[i*cols+j+1] +  
                                              ctop*myData[(i-1)*cols+j])/SPEED;
```

(full code on stencilMPI.c)

MPI Stencil

```
//calculate First row
```

```
if(myId>0){
```

```
    request[1].wait(status);
```

```
    for(int j=1; j<cols-1;j++){
```

```
        buff[j] = myData[j]+( cbotton* myData[cols+j]+
```

```
                                cwest*myData[j-1] +
```

```
                                ceast*myData[j+1] +
```

```
                                ctop*prevRow[j])/SPEED;
```

```
    }
```

MPI Stencil

```
//calculate last row
```

```
if(myId>numP-1){
```

```
    request[3].wait(status);
```

```
    for(int j=1; j<cols-1;j++){
```

```
        buff[(myRows-1)*cols+j]=
```

```
        cbotton*nextRow[j]+
```

```
        cwest*myData[(myRows-1)*cols+j-1]+
```

```
        ceast*myData[(myRows-1)*cols+j+1]+
```

```
        ctop*myData[(myRows-2)*cols+j];
```

```
    }
```

```
//At the end of stencil swap buffers
```

```
Swap(buff with myData and so the for loop again)
```

OpenMP Acceleration Profiling

CPU: Intel i7-7820HQ 8 cores @2.90GHz

gcc -version: 9.3.0

MPI version: OpenMPI 1.2

Compile : mpic++ -g -Wall -o3 -o stencil stencilMPI.c

Run: mpiexec -n 4

Use

a. MPI_Wtime() to time areas in your code

b. Use [gprof](#) or [perf](#) (profilers) to analyze timings and bottleneck on code

Elapsed time:


Non-Blocking: 294 sec

(for an array of size 2^{11} by 2^{11} , 6553 iterations(convolutions) and speed 10^8)

A ~4 Speedup

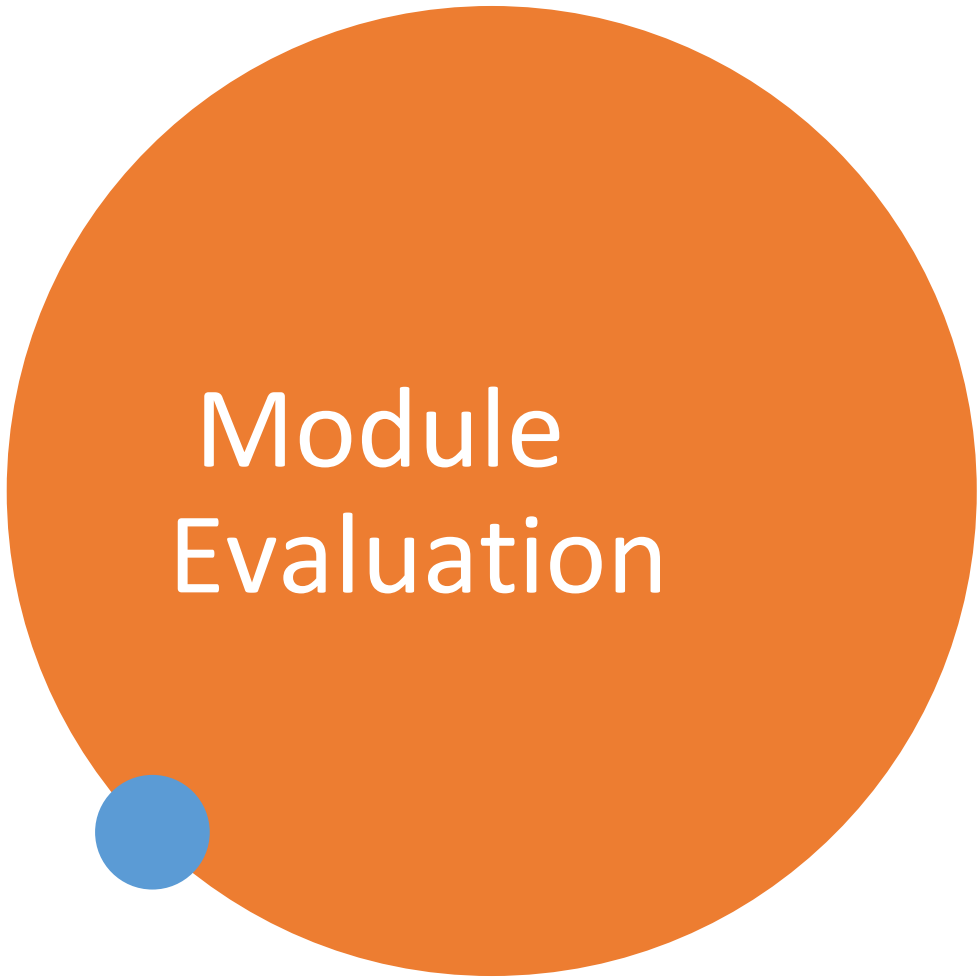


Comparison Sequential vs MPI



Conclusion

- Stencil Code is ubiquitous in Science:
 - Differential Equations Solver
 - System of Equations Solver
 - Heat Transfer
 - Artificial intelligence
 - Etc
- With small changes to original sequential code we achieved speeds up of $\sim 4x$ on an "oldish and cheap" laptop, clearly only running on the multicores of the machines not on a real distributed system



Module Evaluation

- Possible Advance Questions:
 - (code) The stencil code :
Rewrite code for stencil so instead of hardcoding in the loop the neighbor cells. Make the code more flexible and input the neighborhood and its coefficients as a filter. Pass the filter as an input to the stencil function
 - (code) Change it to 3D stencil code
 - (code) Chunk not just by rows do it also by columns