

MPI Collective Communication

Getting started

Login:

```
$ ssh <username>@bw.ncsa.illinois.edu <ENTER>
```

Interactive node request:

```
$ qsub -l -l -l nodes=4:ppn=32:xe,walltime=03:00:00 <ENTER>
```

Download code:

```
$ wget http://shodor.org/~mludin/BW_Capstone/mpi_collective_comm.tar <ENTER>
```

Extract the Zip File:

```
$ tar -xvzf mpi_collective_comm.tar <ENTER>
```

Change folders:

```
$ cd mpi_collective_comm/ <ENTER>
```

```
$ ls -l
```


Message Passing Interface (MPI)

- Standard for distributed memory parallelism.
- Allows for multiple nodes (or just multiple cores) to run a program in parallel.
- Utilizes function calls as opposed to compiler directives.
- Syntax example: send a message:


```
MPI_Send(&buffer, count, MPI_INT, destination, tag, MPI_COMM_WORLD);
```

Distributed Memory Multi-node System

\$ numactl -- hardware




Shared-Memory Threads: (Review)



Shared-Memory: Threads (th0, th1) within a process accessing data.

Distributed-Memory Processes:


Node0:



Distributed-Memory: multiple processes within SPMD accessing data.


Communicator (grouping processes)

- A collection of MPI processes that can send and receive messages to and from each other.
- Normally this is all of the processes, and there is a constant defined for it, `MPI_COMM_WORLD`




Rank

- Unique identifier for each process in the communicator.
- Usually an integer starting at 0 and counting upwards.

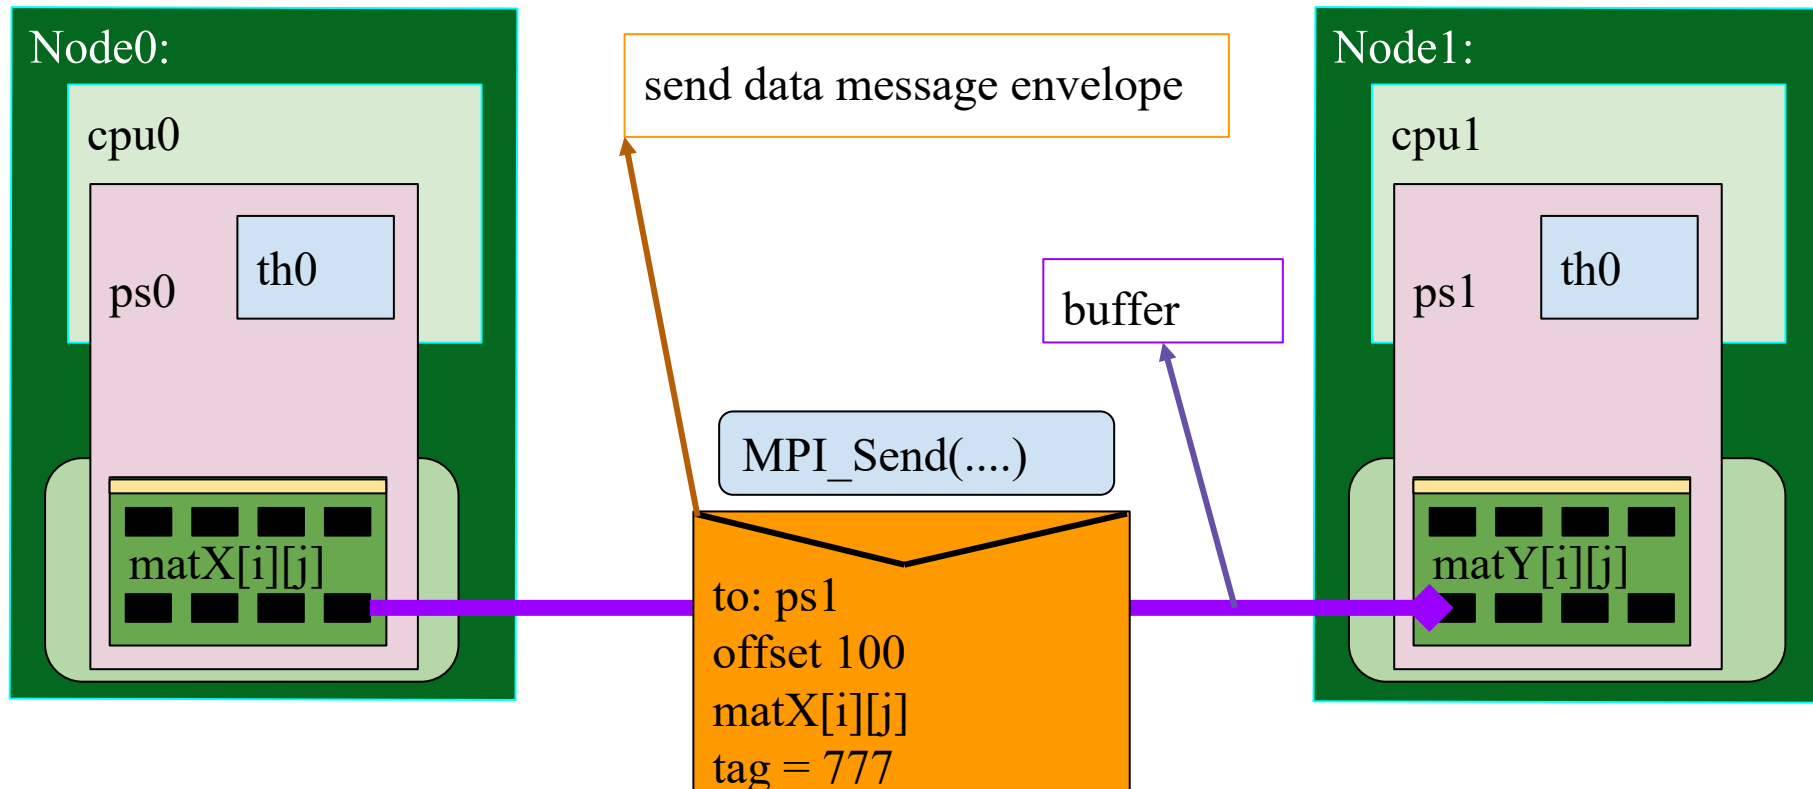


Size

- Number of processes in a communicator.
- Same for all processes in the communicator.

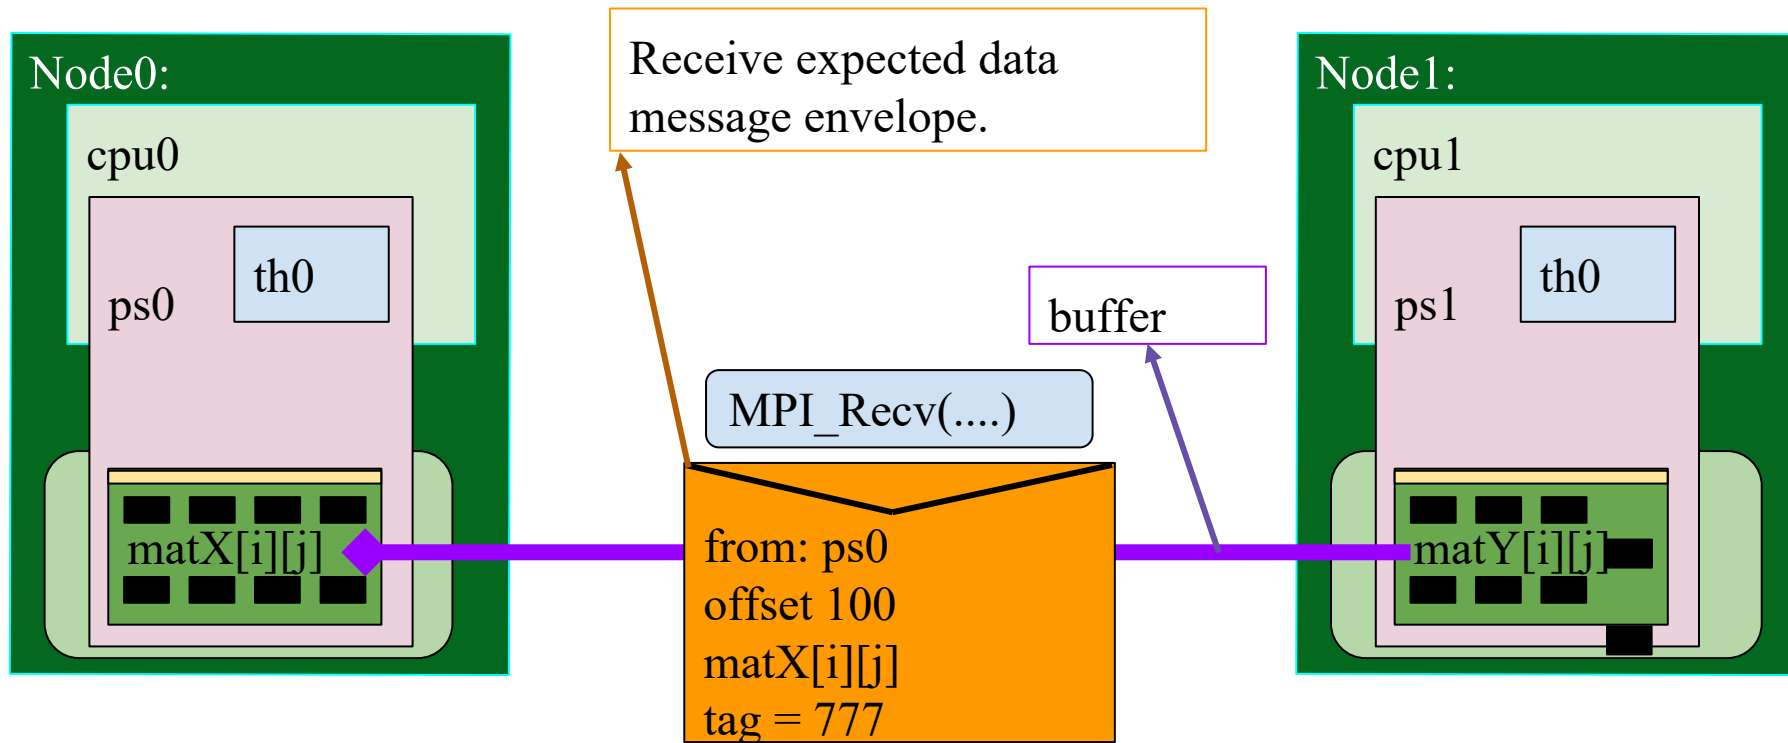


MPI_SEND(...):



ps0 places matrix matX[100-199] from local memory into buffer and calls send routine.

MPI_Recv(...):



ps1 places MPI_Recv(), and awaits until the data from ps0 gets to its buffer, before copy it to local storage

MPI Program Structure:

```
#include <mpi.h>           //include mpi header file.
char message[200];         //message size
int my_rank, num_ps;       //Variable declaration
MPI_Init(&argc, &argv);    // Start MPI Environment now

if (my_rank == master ) {
    MPI_Recv(message)
}
else {
    MPI_Send(&message)
}
MPI_Finalize() //close MPI communication
```

Examples:

[mpi_bcast.c]

```
$ less mpi_bcast.c
```

How to compile:

```
$ make mpi_bcast
```

How to run:

```
$ aprun -n 4 ./mpi_bcast.exe
```

[mpi_reduce.c]

```
$ less mpi_reduce.c
```

How to compile:

```
$ make mpi_reduce
```

How to run:

```
$ aprun -n 4 ./mpi_reduce.exe
```

[mpi_pi_area.c]

```
$ less mpi_pi_area.c
```

How to compile:

```
$ make mpi_pi_area
```

How to run:

```
$ aprun -n 4 ./mpi_pi_area.exe
```

[mpi_scatter.c]

```
$ less mpi_scatter.c
```

How to compile:

```
$ make mpi_scatter
```

How to run:

```
$ aprun -n 4 ./mpi_scatter.exe
```

[mpi_gather.c]

```
$ less mpi_gather.c
```

How to compile:

```
$ make mpi_gather
```

How to run:

```
$ aprun -n 4 ./mpi_gather.exe
```

[mpi_allgather.c]

```
$ less mpi_allgather.c
```

How to compile:

```
$ make mpi_allgather
```

How to run:

```
$ aprun -n 4 ./mpi_allgather.exe
```

References / Further Readings

- a. [Open MPI Organization/Community](#)
- b. [Top500 List](#) of supercomputers in the world
- c. [MPI Library Man Pages](#)
- d. [MPI Standards](#)
- e. [MPI Tutorials](#)
- f. [More MPI Tutorials](#)