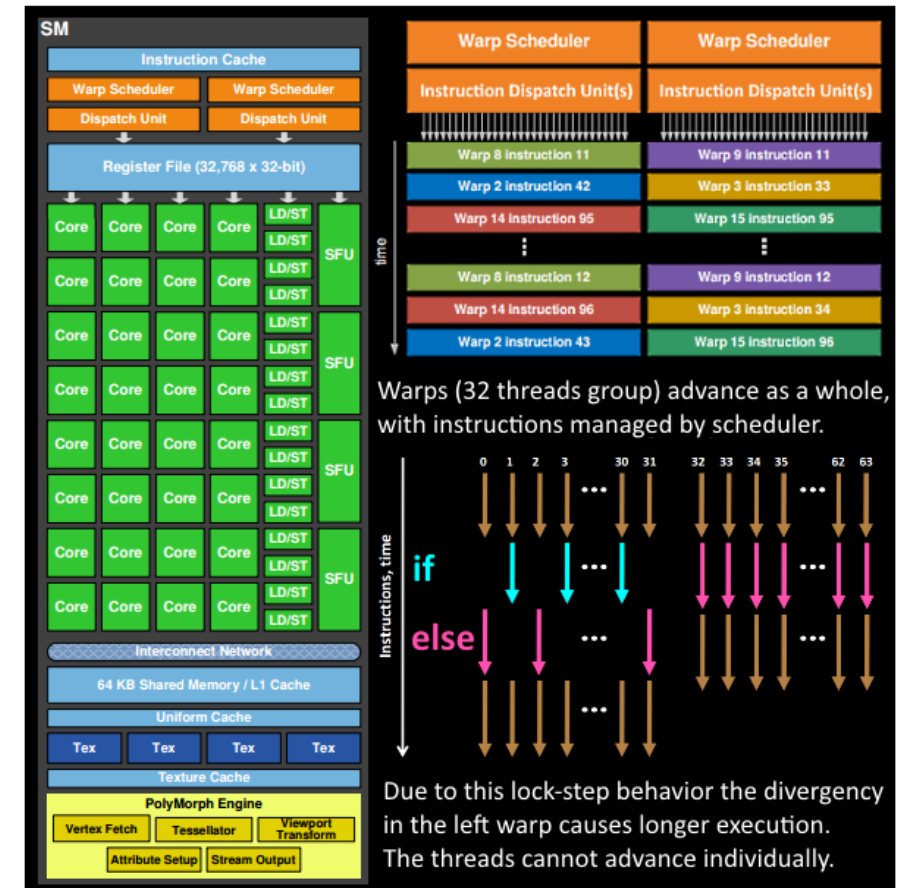


Branching and GPU Efficiency



Many-Core, but not independent

- GPUs are many core, but the structure of memory in the GPU and how threads are organized to access the memory is strictly tied together.
- Threads are bundled in “warps” which have access to the same register and have a bundled set of instructions



Branching can reduce efficiency

- Any time one of the (typically) 32 threads in a warp does something different from the others, all of the others have to wait for the thread to complete its branched activity



What branching looks like in your kernel

```
__global__ void branching(int n)
{
    int threadMod = threadIdx.x%threadsPerBlock;

    if ( threadMod < n) {
        ...
    } else if(threadMod<2*n) {
        ...
    } else if(threadMod<3*n) {
        ...
    } else if(threadMod<4*n) {
        ...
    }
}
```