

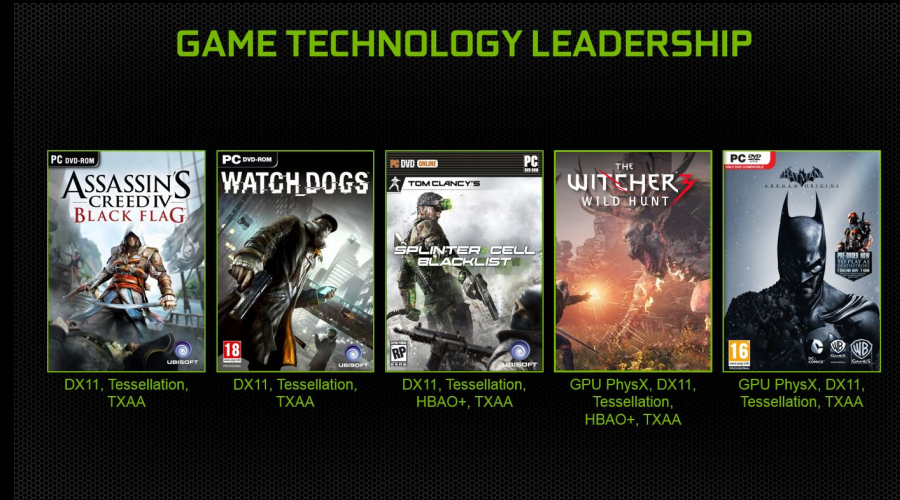
Attribution/License

- Original Materials developed by Mike Shah, Ph.D. (www.mshah.io)
- This work was initially developed for the 2020 PetaScale Blue Waters Institute
- Funding for the development of this work came from <http://shodor.org/>
- This slideset and associated source code may be used freely
 - Attribution is appreciated but not necessary

Introduction to CUDA

Graphics Programming Unit (GPU)

- The Graphics Programming Unit (GPU) is something many of you may be familiar with:
 - GPUs are frequently used to power games and improve their graphics!
 - But modern GPU technology has allowed us to program more than just games
 - GPUs can be used for more general-purpose applications as well!



<https://cdn.wccftech.com/wp-content/uploads/2013/08/NVIDIA-Game-Technology-Leadership.png>

General Purpose Graphics Programming Unit - GPGPU (1/2)

- General Purpose Graphics Programming Unit
 - This is the term given to a more rich programming interface on GPUs
- NVIDIA's CUDA API is one such Application Programming Interface (API) that allows us to write C programs on our GPU
 - https://en.wikipedia.org/wiki/CUDA#Programming_abilities
 - Before CUDA, we were very limited in how expressive of programs we could write on GPUs--again, primarily applications that handled graphics workloads



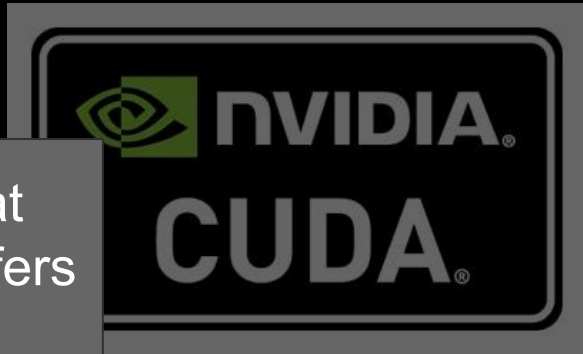
https://upload.wikimedia.org/wikipedia/en/b/b9/Nvidia_CUDA_Logo.jpg

General Purpose Graphics Programming Unit - GPGPU (2/2)

- General Purpose Graphics Programming Unit
 - This is the term given to a more rich programming interface on GPUs

- NVIDIA's CUDA API Programming Interface
to write C programs
- Let's go ahead and look at how a CUDA program differs from C programs that you have been writing

- <https://en.wikipedia.org/wiki/CUDA>
- Before CUDA, we were very limited in how expressive of programs we could write on GPUs--again, primarily applications that handled graphics workloads



upload.wikimedia.org/wikipedia/en/b/b9/Nvidia_CUDA_logo.jpg

A CPU Use Case and a GPU Use Case

CPU Program

- To understand more about how a GPU works, let's first look at solving a problem on the CPU
- We will then solve that exact same problem using the GPU

Problem: Add 1 to every item in an array (1/2)

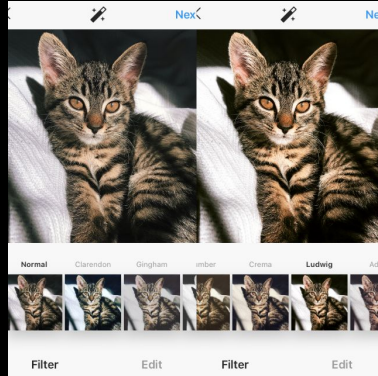
- The program we are going to write is a simple one
- We are going to 'add 1' to every element in an array

Before						
myArray	76	57	42	89	...	7
Index	0	1	2	3	...	n

After						
myArray	77	58	43	90	...	8
Index	0	1	2	3	...	n

Problem: Add 1 to every item in an array (2/2)

- While this is a simple program--it demonstrates a real world use case for something you might do in image processing--such as 'brighten' the pixels in an image



Before						
myArray	76	57	42	89	...	7
Index	0	1	2	3	...	n

After						
myArray	77	58	43	90	...	8
Index	0	1	2	3	...	n

cpuAddOne (in file hello_cpu.c) (1/2)

- Provided to the right, is a solution to this problem
 - We can write a function which iterates through every single element in the array, incrementing the value by 1
- Question:
 - What is the complexity of this algorithm?

```
16 // Increments all values in array by 1
17 void cpuAddOne(int* array,int size){
18     for(unsigned int i=0; i < size; i++){
19         array[i]+=1;
20     }
21 }
```

cpuAddOne (in file hello_cpu.c) (2/2)

- Provided to the right, is a solution to this problem
 - We can write a function which iterates through every single element in the array, incrementing the value by 1
- Question:
 - What is the complexity of this algorithm?
 - $O(n)$ -- because we access every element of the array one time

```
16 // Increments all values in array by 1
17 void cpuAddOne(int* array,int size){
18     for(unsigned int i=0; i < size; i++){
19         array[i]+=1;
20     }
21 }
```

Complete CPU Program

- The entire CPU program is listed to the right
 - We start from our main() function
 - We allocate memory for our array
 - We initialize values in our array to zero
 - We call our add function
 - Then we print our result
- We can fairly easily trace through the execution of this program.

```
1 // Compile with: gcc hello_cpu.cu -o hello_cpu
2 // Run with    : ./hello_cpu
3 // Include our C Standard Libraries
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 // This is our problem size
8 #define SIZE (640*480)
9
10 // Increments all values in array by 1
11 void cpuAddOne(int* array,int size){
12     for(unsigned int i=0; i < size; i++){
13         array[i]+=1;
14     }
15 }
16
17 int main(){
18     // ===== CPU Code =====
19     // Create an array of numbers
20     int* myCPUArray = (int*)malloc(SIZE*sizeof(int));
21     // Let's initialize our values to a specific value to start.
22     // We can do so using an array.
23     for(int i=0; i < SIZE; i++){
24         myCPUArray[i] = 0;
25     }
26     // (For bonus points)
27     // Alternatively you can set each byte to a specific value (see 'memset')
28
29     // Now let's actually solve a problem by incrementing our array.
30     cpuAddOne(myCPUArray,SIZE);
31
32     for(int i=0; i < SIZE; i++){
33         printf("myCPUArray[%i]=%i\n",i,myCPUArray[i]);
34     }
35
36     free(myCPUArray);
37
38     return 0;
39 }
```

Complete CPU Program

- The entire CPU program is listed to the right
 - We start from our main function
 - We allocate memory for our array
 - We initialize values to zero
 - We call our add function
 - Then we print our result
- We can fairly easily trace through the execution of this program.

Let's now take a look at CUDA, and run our first General Purpose GPU Program

```
1 // Compile with: gcc hello_cpu.cu -o hello_cpu
2 // Run with    : ./hello_cpu
3 // Include our C Standard Libraries
4 #include <stdio.h>
5 #include <stdlib.h>
6
7 // This is our problem size
8 #define SIZE (640*480)
9
10 // Increments all values in array by 1
11 void cpuAddOne(int* array,int size){
12     for(unsigned int i=0; i < size; i++){
13         array[i]+=1;
14     }
15 }
16
17 int main(){
18     // Allocate memory for our array
19     int* myCPUArray = (int*) malloc(SIZE*sizeof(int));
20     // Initialize values to zero
21     memset(myCPUArray,0,SIZE*sizeof(int));
22
23     // Call our add function
24     cpuAddOne(myCPUArray,SIZE);
25
26     // (For bonus points)
27     // Alternatively you can set each byte to a specific value (see 'memset')
28
29     // Now let's actually solve a problem by incrementing our array.
30     cpuAddOne(myCPUArray,SIZE);
31
32     for(int i=0; i < SIZE; i++){
33         printf("myCPUArray[%i]=%i\n",i,myCPUArray[i]);
34     }
35
36     free(myCPUArray);
37
38     return 0;
39 }
```

GPU Program (in file hello_gpu.cu) (1/2)

- To the right is the GPU code that will do the exact same thing as before!
- Don't worry, we will walk through the code in more detail.

```
63 __global__ void gpuAddOne(int* array, int size){
64     // We need some unique variable identifying which thread accesses
65     // a piece of data. So below we are going to figure out which
66     // thread (i.e. index) is going to operate on which piece of data.
67     //
68     // This gives us full coverage of every id that we want to access.
69     // Figure out which block we are, what dimension, and which thread.
70     // You can think of these values like a phone number if you like.
71     // The block is the area code for example.
72     int currentThread = (blockIdx.x * blockDim.x) + threadIdx.x;
73     // Do a bounds check to make sure we are not accessing
74     // out of bounds memory.
75     // (This can occur if we don't have a perfect number within 32.
76     if(currentThread < size){
77         array[currentThread] += 1;
78     }
79 }
```

GPU Program (in file hello_

- Here is a full view of the code, so we can see a full sample

```
1 // Compile with: nvcc hello_gpu.cu -o hello_gpu
2 // Run with : ./hello_gpu
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 // This is our problem size
7 #define SIZE (640*480)
8
9 __global__ void gpuAddOne(int* array, int size){
10     int currentThread = (blockIdx.x * blockDim.x) + threadIdx.x;
11     if(currentThread < size){
12         array[currentThread] += 1;
13     }
14 }
15 int main(){
16     // ===== CPU Code =====
17     // Create an array of numbers
18     int* myCPUArray = (int*)malloc(SIZE*sizeof(int));
19     // Let's initialize our values to a specific value to start.
20     // We can do so using an array.
21     for(int i=0; i < SIZE; i++){
22         myCPUArray[i] = 0;
23     }
24     // ===== CPU Code =====
25     // ===== GPU Code =====
26     int* myGPUArray = (int*)malloc(SIZE*sizeof(int));
27     cudaMalloc(&myGPUArray, SIZE*sizeof(int));
28     cudaMemcpy(myGPUArray, myCPUArray, SIZE*sizeof(int), cudaMemcpyHostToDevice);
29     int NUM_THREADS = 256;
30     int NUM_BLOCKS = (int)ceil(SIZE/ (float)NUM_THREADS);
31     gpuAddOne<<<NUM_BLOCKS,NUM_THREADS>>>(myGPUArray,SIZE);
32     cudaMemcpy(myCPUArray, myGPUArray, SIZE*sizeof(int), cudaMemcpyDeviceToHost);
33
34     for(int i=0; i < SIZE; i++){
35         printf("myCPUArray [%i]=%i\n",i,myCPUArray[i]);
36     }
37
38     cudaFree(myGPUArray);
39     // ===== GPU Code =====
40     // Free our CPU Memory
41     free(myCPUArray);
42     return 0;
43 }
```

GPU Program (in file hello_gpu.cu) (2/2)

- Question to audience:
 - What is the $O(n)$ of this code?
 - ?!

```
63 __global__ void gpuAddOne(int* array, int size){
64     // We need some unique variable identifying which thread accesses
65     // a piece of data. So below we are going to figure out which
66     // thread (i.e. index) is going to operate on which piece of data.
67     //
68     // This gives us full coverage of every id that we want to access.
69     // Figure out which block we are, what dimension, and which thread.
70     // You can think of these values like a phone number if you like.
71     // The block is the area code for example.
72     int currentThread = (blockIdx.x * blockDim.x) + threadIdx.x;
73     // Do a bounds check to make sure we are not accessing
74     // out of bounds memory.
75     // (This can occur if we don't have a perfect number within 32.
76     if(currentThread < size){
77         array[currentThread] += 1;
78     }
79 }
```


GPU Program (in file hello_gpu.cu) (2/2)

- Question to audience:
 - What is the $O(n)$ of this code?
 - Answer is closer to $O(1)$ -- and we'll need to look at our GPU to understand why!

```
63 __global__ void gpuAddOne(int* array, int size){
64     // We need some unique variable identifying which thread accesses
65     // a piece of data. So below we are going to figure out which
66     // thread (i.e. index) is going to operate on which piece of data.
67     //
68     // This gives us full coverage of every id that we want to access.
69     // Figure out which block we are, what dimension, and which thread.
70     // You can think of these values like a phone number if you like.
71     // The block is the area code for example.
72     int currentThread = (blockIdx.x * blockDim.x) + threadIdx.x;
73     // Do a bounds check to make sure we are not accessing
74     // out of bounds memory.
75     // (This can occur if we don't have a perfect number within 32.
76     if(currentThread < size){
77         array[currentThread] += 1;
78     }
79 }
```

GPU Program (in file hello_gpu.cu) (2/2)

- Question to audience:

- What is the $O(n)$ of this code?

- Answer is close to $O(n)$ and we'll need our GPU to understand why!

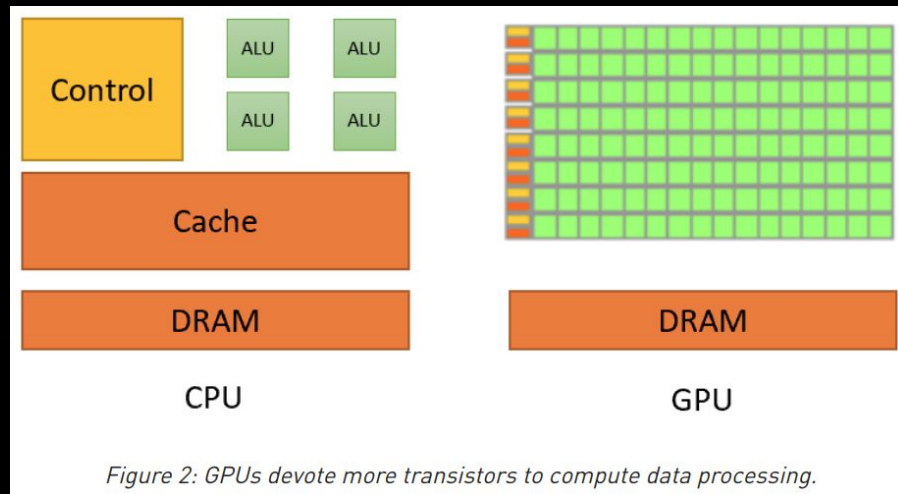
Let's understand a bit more about 'why'

```
63 __global__ void gpuAddOne(int* array, int size){
64     // We need some unique variable identifying which thread accesses
65     // a piece of data. So below we are going to figure out which
66     // thread (i.e. index) is going to operate on which piece of data.
67     //
68     // We need to figure out the range of every id that we want to access.
69     // We are, what dimension, and which thread.
70     // values like a phone number if you like.
71     // code for example.
72     int Idx.x * blockDim.x) + threadIdx.x;
73     // make sure we are not accessing
74     //
75     // don't have a perfect number within 32.
76     //
77     // = 1;
78     //
79 }
```

GPU Architecture

GPU Architecture (1/2)

- The first thing we need to understand in our use case, is that one code is written on the CPU, and one for the GPU
- GPU's (right) are structured differently than our CPUs (left)
- Notice the little 'green' boxes in the GPU--there are many more of these compute units--and those units are meant to execute code in parallel in individual **threads**.



<https://developer.nvidia.com/blog/cuda-refresher-reviewing-the-origins-of-gpu-computing/>

GPU Architecture (2/2)

- The first thing we need to understand in our use case, is that one code is written and one for the GPU
- GPU's (right) are structured differently than our
- Notice the little 'green' units in the GPU--there are many more of these compute units--and those units are meant to execute code in parallel in individual **threads**.

So what is a thread, and how do we execute code in these 'boxes'

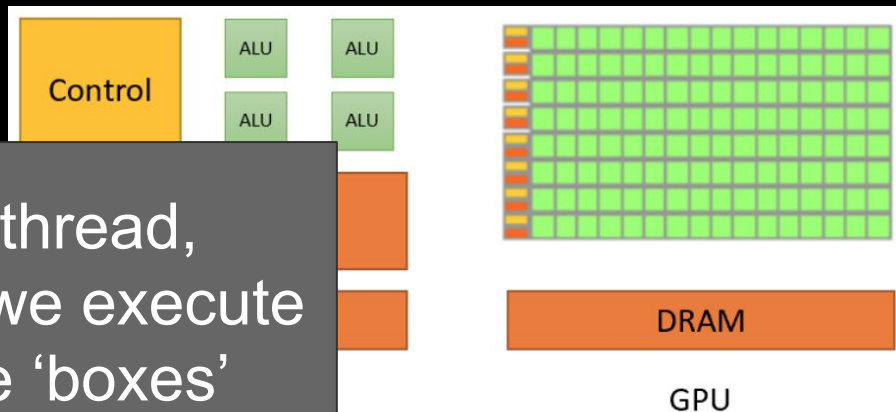


Figure 2: GPUs devote more transistors to compute data processing.

<https://developer.nvidia.com/blog/cuda-refresher-reviewing-the-origins-of-gpu-computing/>

Threads

- Threads are traditionally referred to as 'lightweight' processes
- They are used to execute a part of your program concurrently while your program runs.
- In CUDA, we execute many threads (100s or 1000s) in parallel in order to get a speedup in execution.
 - So if we look at our example to the right
 - We can see we a 'single thread' that is going to execute this tiny gpuAddOne program on our GPU

```
63 __global__ void gpuAddOne(int* array, int size){
64     // We need some unique variable identifying which thread accesses
65     // a piece of data. So below we are going to figure out which
66     // thread (i.e. index) is going to operate on which piece of data.
67     //
68     // This gives us full coverage of every id that we want to access.
69     // Figure out which block we are, what dimension, and which thread.
70     // You can think of these values like a phone number if you like.
71     // The block is the area code for example.
72     int currentThread = (blockIdx.x * blockDim.x) + threadIdx.x;
73     // Do a bounds check to make sure we are not accessing
74     // out of bounds memory.
75     // (This can occur if we don't have a perfect number within 32.
76     if(currentThread < size){
77         array[currentThread] += 1;
78     }
79 }
```

Threads Continued - Single Instruction Multiple Thread (SIMT)

- In CUDA we call this model SIMT, for single instruction multiple thread
- We are going to have many (100s or 1000s) of threads executing on each of our GPUs execution units.

```
63 __global__ void gpuAddOne(int* array, int size){
64     // We need some unique variable identifying which thread accesses
65     // a piece of data. So below we are going to figure out which
66     // thread (i.e. index) is going to operate on which piece of data.
67     //
68     // This gives us full coverage of every id that we want to access.
69     // Figure out which block we are, what dimension, and which thread.
70     // You can think of these values like a phone number if you like.
71     // The block is the area code for example.
72     int currentThread = (blockIdx.x * blockDim.x) + threadIdx.x;
73     // Do a bounds check to make sure we are not accessing
74     // out of bounds memory.
75     // (This can occur if we don't have a perfect number within 32.
76     if(currentThread < size){
77         array[currentThread] += 1;
78     }
79 }
```

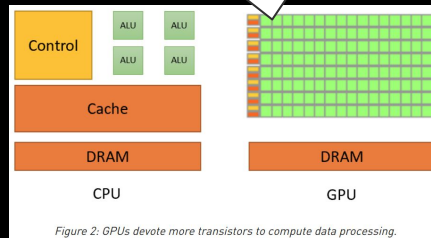


Figure 2: GPUs devote more transistors to compute data processing.

Threads Continued - Single Instruction Multiple Thread (SIMT)

- So in CUDA, the lowest granularity of execution is a 'thread'.
 - A thread executes a series of instructions (i.e. a small program).
- Groups of threads (with consecutive thread indexes) are divided into warps which execute on a single CUDA core.
 - The lowest schedulable entity is known as a 'warp'

```
63 __global__ void gpuAddOne(int* array, int size){
64     // We need some unique variable identifying which thread accesses
65     // a piece of data. So below we are going to figure out which
66     // thread (i.e. index) is going to operate on which piece of data.
67     //
68     // This gives us full coverage of every id that we want to access.
69     // Figure out which block we are, what dimension, and which thread.
70     // You can think of these values like a phone number if you like.
71     // The block is the area code for example.
72     int currentThread = (blockIdx.x * blockDim.x) + threadIdx.x;
73     // Do a bounds check to make sure we are not accessing
74     // out of bounds memory.
75     // (This can occur if we don't have a perfect number within 32.
76     if(currentThread < size){
77         array[currentThread] += 1;
78     }
79 }
```

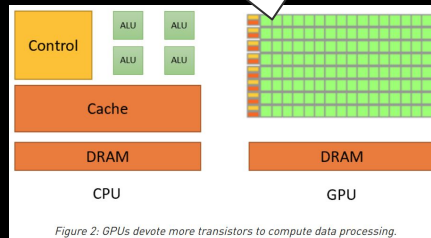


Figure 2: GPUs devote more transistors to compute data processing.

Threads Continued - Single Instruction Multiple Thread (SIMT)

- Thread blocks are the lowest programmable entity.
 - A single shader core is assigned to each thread.
 - They are organized in a 3-D block.
- Thread blocks are then further assigned to a GPU.

```
63 __global__ void gpuAddOne(int* array, int size){
64     // We need some unique variable identifying which thread accesses
65     // a piece of data. So below we are going to figure out which
66     // thread (i.e. index) is going to operate on which piece of data.
67     //
68     // This gives us full coverage of every id that we want to access.
69     // Figure out which block we are, what dimension, and which thread.
70     // You can think of these values like a phone number if you like.
71     // The block is the area code for example.
72     int currentThread = (blockIdx.x * blockDim.x) + threadIdx.x;
73     // Do a bounds check to make sure we are not accessing
74     // out of bounds memory.
75     // (This can occur if we don't have a perfect number within 32.
76     if(currentThread < size){
77         array[currentThread] += 1;
78     }
79 }
```

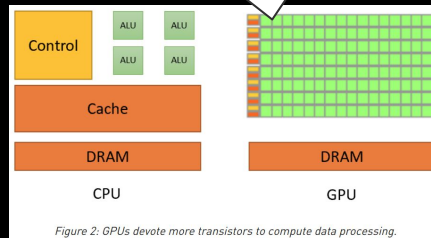


Figure 2: GPUs devote more transistors to compute data processing.

Further Resources

- CUDA Programming Model
 - https://cvw.cac.cornell.edu/GPU/simt_warp