

# **Blue Waters Petascale Semester Curriculum v1.0**

## **Unit 9: Optimization**

### **Lesson 1: Cache Efficient Matrix Multiplication**

*Developed by Paul F. Hemler*

*for the Shodor Education Foundation, Inc.*

Except where otherwise noted, this work by The Shodor Education Foundation, Inc. is licensed under CC BY-NC 4.0. To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc/4.0>

Browse and search the full curriculum at <http://shodor.org/petascale/materials/semester-curriculum>

We welcome your improvements! You can submit your proposed changes to this material and the rest of the curriculum in our GitHub repository at <https://github.com/shodor-education/petascale-semester-curriculum>

We want to hear from you! Please let us know your experiences using this material by sending email to [petascale@shodor.org](mailto:petascale@shodor.org)

# Effective Caching for Matrix Multiplication

# Cache Memory

- ▶ An important part of the memory hierarchy in any computer system
- ▶ Utilizes different technology compared to RAM
  - ▶ Faster
  - ▶ Uses more power
  - ▶ Costs more
- ▶ Gives the illusion of a large (RAM size), fast (cache speed) memory
- ▶ Compromise between cost, access time and size
- ▶ Program execution depends on efficiently utilizing the cache

# Matrix Multiplication

- ▶ Matrices and matrix multiplication are very common in a variety of Engineering and Scientific Computing problems
- ▶ An  $n \times n$  matrix  $A$  is mathematically written as:

$$A = \begin{bmatrix} a_{0,0} & a_{0,1} & \cdots & a_{0,n-1} \\ a_{1,0} & a_{1,1} & \cdots & a_{1,n-1} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n-1,0} & a_{n-1,1} & \cdots & a_{n-1,n-1} \end{bmatrix}$$

- ▶ Matrix multiplication is mathematically written as:

$$C = AB$$

- ▶ Where  $A$ ,  $B$  and  $C$  are all  $n \times n$  matrices

# Matrix Multiplication

- ▶ Each element in  $C$  requires  $n$  multiplications and  $n - 1$  additions
- ▶ Each element in  $C$  is computed by multiplying a row of matrix  $A$  with a column of matrix  $B$ , for example,

$$\begin{bmatrix} c_{0,0} & c_{0,1} & \cdots & c_{0,n-1} \\ c_{1,0} & c_{1,1} & \cdots & c_{1,n-1} \\ \vdots & \vdots & \ddots & \vdots \\ c_{n-1,0} & c_{n-1,1} & \cdots & c_{n-1,n-1} \end{bmatrix} = \begin{bmatrix} a_{0,0} & a_{0,1} & \cdots & a_{0,n-1} \\ a_{1,0} & a_{1,1} & \cdots & a_{1,n-1} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n-1,0} & a_{n-1,1} & \cdots & a_{n-1,n-1} \end{bmatrix} * \begin{bmatrix} b_{0,0} & b_{0,1} & \cdots & b_{0,n-1} \\ b_{1,0} & b_{1,1} & \cdots & b_{1,n-1} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n-1,0} & b_{n-1,1} & \cdots & b_{n-1,n-1} \end{bmatrix}$$

- ▶  $c_{1,0}$  is the sum of the products of the elements of row 1 of matrix  $A$  and column 0 of matrix  $B$
- ▶ Or more generally,

$$c_{i,j} = \sum_{k=0}^{n-1} a_{i,k} b_{k,j}$$

# Matrix Memory Access

- ▶ The main memory of a computer system can be thought of as a long linear array, where elements are stored at consecutive memory locations
- ▶ A matrix is stored in memory either by rows (row major order) or columns (column major order)
- ▶ The **C** language uses row major order, while Fortran uses column major order
- ▶ A matrix in **C** is stored as:



- ▶ Optimal cache utilization occurs when sequential memory elements are accessed
  - ▶ For the matrix multiplication example above, matrix A is efficiently accessed but matrix B is not

# Matrix Multiplication Transpose

- ▶ Element access to matrix  $B$  can be made cache efficient by storing matrix  $B$  in column major order
- ▶ This is the same as computing the *transpose* of  $B$
- ▶ In this case, the elements of the rows of matrix  $A$  are multiplied by the elements of the rows of matrix  $B$

$$\begin{bmatrix} c_{0,0} & c_{0,1} & \cdots & c_{0,n-1} \\ c_{1,0} & c_{1,1} & \cdots & c_{1,n-1} \\ \vdots & \vdots & \ddots & \vdots \\ c_{n-1,0} & c_{n-1,1} & \cdots & c_{n-1,n-1} \end{bmatrix} = \begin{bmatrix} a_{0,0} & a_{0,1} & \cdots & a_{0,n-1} \\ a_{1,0} & a_{1,1} & \cdots & a_{1,n-1} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n-1,0} & a_{n-1,1} & \cdots & a_{n-1,n-1} \end{bmatrix} * \begin{bmatrix} b_{0,0} & b_{1,0} & \cdots & b_{n-1,0} \\ b_{0,1} & b_{1,1} & \cdots & b_{n-1,1} \\ \vdots & \vdots & \ddots & \vdots \\ b_{0,n-1} & b_{1,n-1} & \cdots & b_{n-1,n-1} \end{bmatrix}$$

$$c_{i,j} = \sum_{k=0}^{n-1} a_{i,k} b_{j,k}$$



# Matrix Multiplication

- ▶ Matrix multiplication using the transpose of the  $B$  matrix more efficiently utilizes the cache
- ▶ Both matrix multiplication techniques inefficiently use the cache because the data must be brought into the cache multiple times
  - ▶ Each row of matrix  $A$  is multiplied by each row of matrix  $B$
  - ▶ When the matrices are large, the rows of matrix  $B$  will need to be brought into the cache for each row in matrix  $A$

# Block Matrix Multiplication

- ▶ A better way of utilizing the cache is to perform block matrix multiplication, where the matrices are broken into blocks or tiles
- ▶ For example, a 4 X 4 matrix  $A$  can be made of 2 X 2 blocks or tiles each of size 2 X 2

$$A = \begin{bmatrix} a_{0,0} & a_{0,1} & a_{0,2} & a_{0,3} \\ a_{1,0} & a_{1,1} & a_{1,2} & a_{1,3} \\ a_{2,0} & a_{2,1} & a_{2,2} & a_{2,3} \\ a_{3,0} & a_{3,1} & a_{3,2} & a_{3,3} \end{bmatrix} = \begin{bmatrix} \begin{bmatrix} a_{0,0} & a_{0,1} \end{bmatrix} & \begin{bmatrix} a_{0,2} & a_{0,3} \end{bmatrix} \\ \begin{bmatrix} a_{1,0} & a_{1,1} \end{bmatrix} & \begin{bmatrix} a_{1,2} & a_{1,3} \end{bmatrix} \\ \begin{bmatrix} a_{2,0} & a_{2,1} \end{bmatrix} & \begin{bmatrix} a_{2,2} & a_{2,3} \end{bmatrix} \\ \begin{bmatrix} a_{3,0} & a_{3,1} \end{bmatrix} & \begin{bmatrix} a_{3,2} & a_{3,3} \end{bmatrix} \end{bmatrix} = \begin{bmatrix} A_{0,0} & A_{0,1} \\ A_{1,0} & A_{1,1} \end{bmatrix}$$

- ▶ Block matrix multiplications is

$$\begin{bmatrix} C_{0,0} & C_{0,1} \\ C_{1,0} & C_{1,1} \end{bmatrix} = \begin{bmatrix} A_{0,0} & A_{0,1} \\ A_{1,0} & A_{1,1} \end{bmatrix} \begin{bmatrix} B_{0,0} & B_{0,1} \\ B_{1,0} & B_{1,1} \end{bmatrix}$$

# Block Matrix Multiplication

- The product matrix

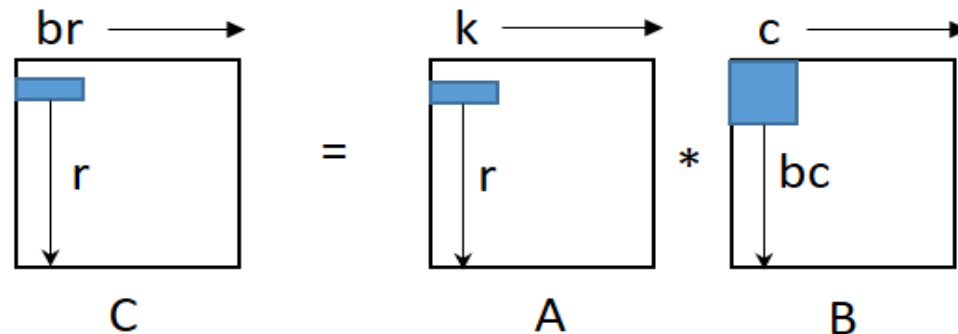
$$C_{0,0} = A_{0,0}B_{0,0} + A_{0,1}B_{1,0}$$

$$C_{0,1} = A_{0,0}B_{0,1} + A_{0,1}B_{1,1}$$

$$C_{1,0} = A_{1,0}B_{0,0} + A_{1,1}B_{1,0}$$

$$C_{1,1} = A_{1,0}B_{0,1} + A_{1,1}B_{1,1}$$

- This technique efficiently utilizes the cache because all the elements in a block in matrix  $B$  are completely used for all computations and are not needed again

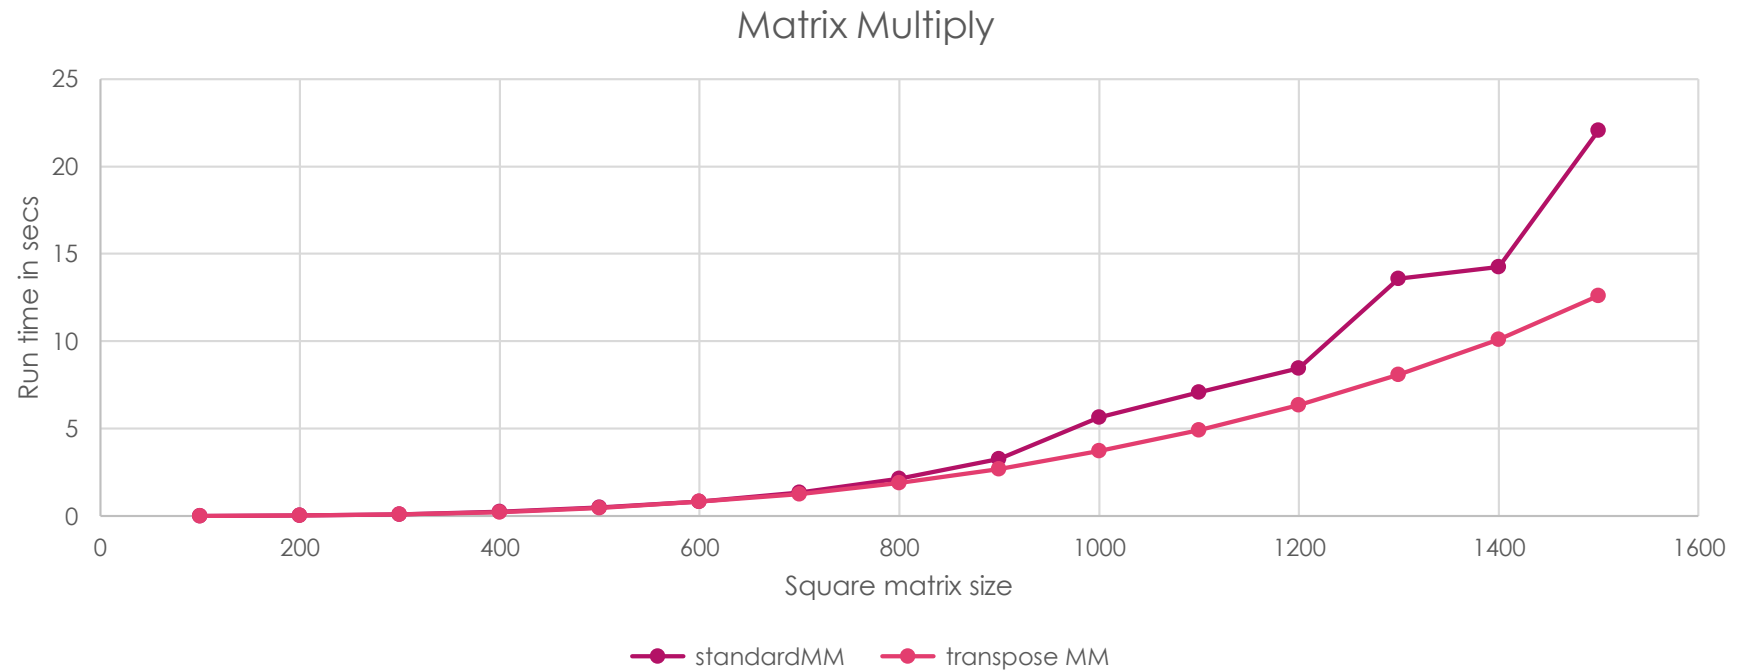


# Matrix Multiply Program

- ▶ Two programs were written to time the standard matrix multiplication with both the transpose and the block matrix techniques
- ▶ Command line arguments are used to determine the size of the matrices
- ▶ All matrices utilize double-precision floating-point numbers
- ▶ Computational results are compared to ensure the product matrix is correctly determined

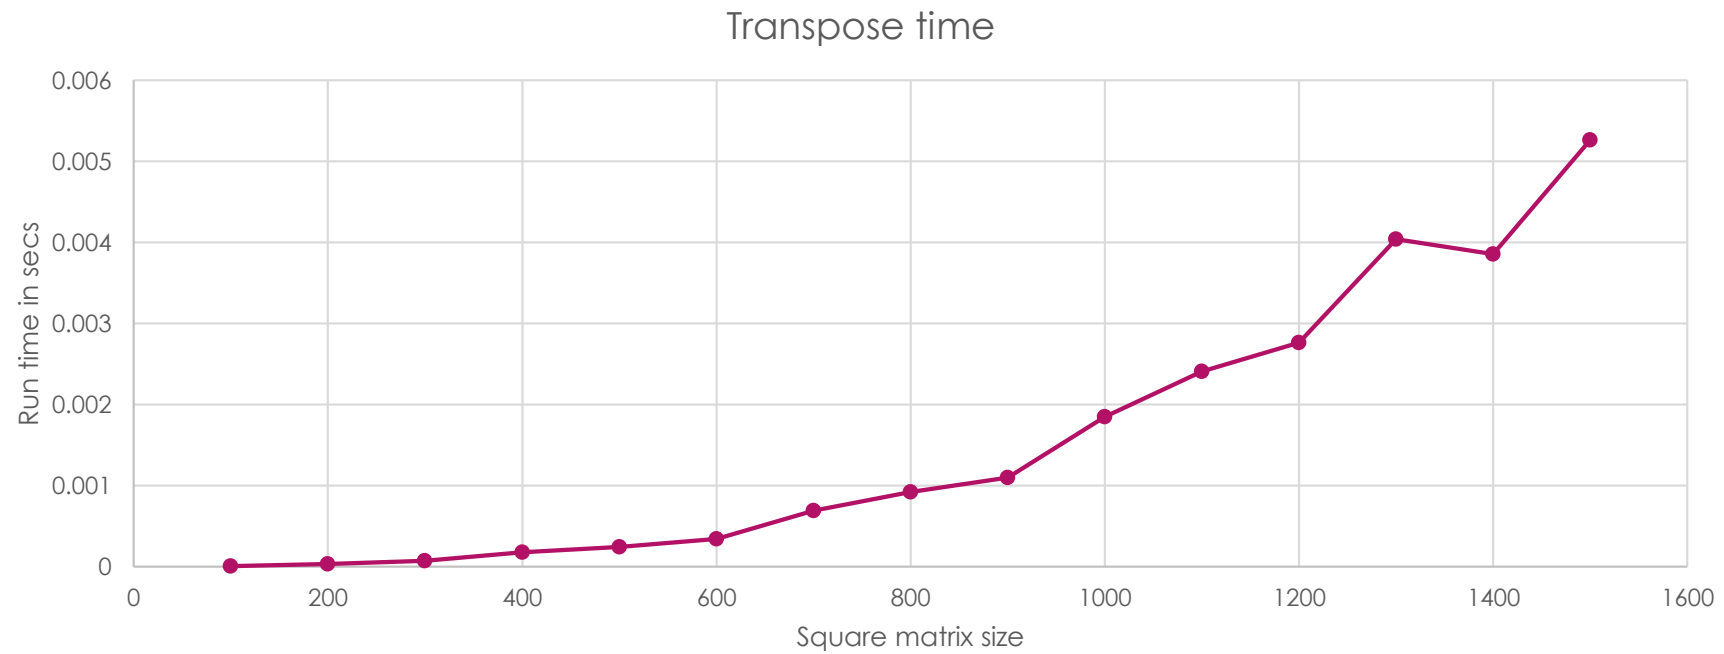
# Program Results

- Timing results for standard matrix multiplication compared to transpose matrix multiplication



# Program Results


- The time to transpose the matrix is negligible



100

- The time to perform block matrix multiplication depends on the size of the block
- Blocks of size 8 X 8 or 16 X 16 appear to give the best performance

## Block Matrix Multiply with doubles



# Program Results

- ▶ Block matrix multiply with blocks of size 16 X 16 appear to give the best performance

